



Title	Standards That Can Lose: An Epistemic-Fiduciary Constitution for the University
Author	Peter Kahl · Director, Lex et Ratio Ltd, United Kingdom · ORCID 0009-0003-1616-4843
Reference	LXR-2026-LAW-EPISTO-WP · v3.0 · 18 July 2026 · doi:10.13140/RG.2.2.14076.09603 and doi:10.5281/zenodo.21468283
Series	Lex et Ratio Working Papers – Fiduciary and public-law governance strand
Type / Status	Working paper, prepared for submission to a law journal. Theoretical and normative. Not peer reviewed. Self-deposited, open access. Approx. 22,500 words.
Companion to	Peter Kahl, 'The Recognition Game: Metric Governance and the Stabilisation of Epistemic Clientelism' (version 2.0, Lex et Ratio Ltd 2026) doi:10.5281/zenodo.20774711 – which supplies the diagnosis of recognition scarcity, epistemic clientelism, evaluative-layer capping and the adversarial-clientelist equilibrium on which Chapters 2 to 7 of this paper draw. That paper asks why the degraded order is stable; this one asks what would have to be true of an institution for the order to be capable of losing.
Jurisdiction	England and Wales for the fiduciary doctrine (<i>Mothew; Keech; Rukhadze</i>), the higher education complaints jurisdiction (Higher Education Act 2004, s 12), and the data protection analysis at §5.2 (UK GDPR and Data Protection Act 2018, as amended by the Data (Use and Access) Act 2025). The argument is jurisdiction-neutral: the mechanisms are stated so as to be portable to any system possessing a sector regulator not itself constituted by the regulated institutions, and §5.4 states what to do where none exists.
Fields	Fiduciary law and theory; administrative and public law; institutional accountability; data protection and research ethics; social epistemology and epistemic injustice; research assessment and audit studies; higher education governance and the philosophy of higher education (as the paper's worked domain).
Keywords	fiduciary law; epistemic fiduciary; fiduciary answerability; duty of loyalty; fiduciary breach; delegated discretion; fettering of discretion; standard attribution; standing to trigger; epistemic finality; representational sealing; hollowed authorisation; epistemic capture; epistemocracy; periodic review against later evidence; epistemic justice; contributory injustice; epistemic clientelism; evaluative-layer capping; metric governance; audit society; optocratic drift; fiducial hollowing; epistemic inversion; higher education governance
Audience	Scholars of fiduciary and public law, and of social epistemology; those designing oversight, standing and appeal mechanisms in any sector where an authority controls whether its own determinations are reopened; sector regulators, funders and research-assessment bodies; university governing bodies, councils and senates; and student and staff representative bodies – who are, on this paper's account, the parties to whom the duty is owed rather than parties to be consulted.
Related Works	Three results are imported and stated without re-argument (§1.2): Peter Kahl, 'The Inside of Generatedness' (2026) doi:10.5281/zenodo.20353457 (the seal); 'Answerable Knowledge' (2026) doi:10.5281/zenodo.20439758 (the operative conditions); and 'Epistemic Humility as Fiduciary Obligation' (2026) doi:10.5281/zenodo.18440605 (the beneficiary). The fourth import, 'The Recognition Game', is the <i>Companion</i> to work above. Diagnostic predecessor: 'The Fiduciary Breach of the Modern University' (2025) doi:10.5281/zenodo.17618878 (§§1.1, 3.1, 5.12). Also drawn upon: 'Conformance Is Not Constitution' (2026) doi:10.13140/RG.2.2.28710.28486 and 'Delegated Discretion' (2026) doi:10.5281/zenodo.20920750 (§§1.1, 5.13); 'Out-of-Distribution Failure' (2026) doi:10.5281/zenodo.19353400 (§5.2); and, earlier, 'Directors' Epistemic Duties and Fiduciary Openness' (2nd edn, 2025) doi:10.13140/RG.2.2.28155.17449 and 'Epistemic Justice and Institutional Responsibility in Academia' (2nd edn, 2025) doi:10.13140/RG.2.2.11063.38563.
Declaration of interests	None. The work is unfunded. The author is sole director of Lex et Ratio Ltd, the publisher of this series, and holds no post in, funding from, or contractual relationship with any higher education institution. The paper is self-deposited and not peer reviewed; its author is, on its own argument, the party least able to certify it, and the conditions under which it is defeated are stated at §6.5 rather than left to be supplied by a reader.
Licence	Creative Commons Attribution 4.0 International (CC BY 4.0)
Disclaimer	Published for academic and policy-discussion purposes only; not legal, regulatory, financial or technical advice. Views are the author's own.
Cite as	Peter Kahl, 'Standards That Can Lose: An Epistemic-Fiduciary Constitution for the University' (Lex et Ratio Ltd Working Paper LXR-2026-LAW-EPISTO-WP, Version 3.0, 18 July 2026) doi:10.13140/RG.2.2.14076.09603, and doi:10.5281/zenodo.21468283.

Abstract

This paper proposes that the university be constituted as an epistemic fiduciary in respect of its evaluative powers — the powers by which it admits, assesses, credits, promotes, funds and recognises — and that the defining incident of that constitution be a single structural requirement: that each such standard remain *capable of losing*, defeasible by what happens to those it governs, through a route the standard's holder does not control. It argues not that universities are fiduciaries of knowledge as a matter of positive law — they are not — but that they should be so constituted, and it works out what that constitution requires. The resulting model is named *epistemocracy*, in a sense deliberately opposed to the word's established grain: not rule by those certified as knowing, but governance in which epistemic authority is institutionally answerable and no recognised knower holds finality. The model is grounded in fiduciary answerability, epistemic plurality, and distributed credibility. Building on fiduciary theory (Frankel, Smith, Miller), virtue epistemology (Fricker, Zagzebski, Medina), and critical pedagogy (Freire, Darder), I diagnose the epistemic crisis of the contemporary university through three critical terms — *optocratic drift*, *fiducial hollowing*, and *epistemic inversion* — which identify how institutions privilege optics over substance, hollow fiduciary duties into public relations, and collapse plurality into singular figureheads.

The fiduciary claim is advanced as a proposal for how the university ought to be constituted, not as a description of English law: fiduciary status follows an undertaking of loyalty (*Bristol and West Building Society v Mothew*), and no such undertaking presently runs from the university to those its evaluative standards govern. Against this backdrop I propose epistemocracy as a structural reform whose primary mechanism is the periodic review of recognition standards against later evidence — admitted through a three-tier evidential architecture, in which later success among the excluded triggers review, a class-level pattern shifts the burden, and defeat requires that pattern to exceed a pre-specified error rate — supported by oversight bodies that the reviewed authority does not appoint, resource, instruct or dissolve, and which any affected party has standing to compel. The proposal is set against the oversight architecture the sector actually builds, which reproduces the pathologies it is meant to answer: bodies no beneficiary has standing to invoke, appointed by the institution they review, defended by the plurality of their composition rather than by independence from what they examine, and justified by their capacity to produce legitimacy — the defect Michael Power identifies in audit generally.

The model rests instead on exposure: the requirement that an evaluative standard remain capable of *losing*, by a route its holder does not control. Two consequences follow, and both are uncomfortable. Transparency, disclosure and reason-giving, though necessary and valuable, cannot by themselves supply the missing check, because a standard's holder cannot exhaustively and independently validate the standard under which their own representing proceeds. And a body performing the function proposed here will frequently *reduce* institutional legitimacy rather than restore it — which is the point, and the test of whether it is working. The paper specifies what follows when a standard loses, and provides that where the operative standard was constituted outside the institution applying it — by a research-assessment framework, accreditor or funder — the finding must travel to that constitutor rather than terminate at the point of application. Beyond academia, epistemocracy provides a transferable framework for corporations, NGOs and democratic institutions, on conditions stated in Chapter 8.

About the Author

Peter Kahl is an independent researcher at Lex et Ratio Ltd whose work examines institutional accountability, evaluative governance and the constitutional architecture of knowledge institutions. He holds an LLB and an LLM, and his research explores how standards are constituted, recognised and revised across law, artificial intelligence and higher education. A recurring theme of his work is the distinction between exercising authority under an evaluative framework and possessing the authority to justify or reconstruct that framework itself.

His recent scholarship argues that many contemporary institutions, including universities, increasingly govern through recognition systems that determine what counts as knowledge, merit and innovation. His work therefore focuses not only on the decisions institutions make, but on the evaluative structures that shape which questions may be asked, which contributions become visible, and which forms of knowledge remain unrecognised.

Before entering legal research, he spent more than two decades as an electronics engineer, including analog and mixed-signal integrated circuit design at National Semiconductor in Silicon Valley. That engineering background informs his approach to institutional analysis: robust systems require not only correct operation, but traceability, reconstructability and the ability to identify who is answerable for the standards under which they operate.

About Lex et Ratio Ltd

Lex et Ratio Ltd is an independent legal research and policy practice specialising in evaluative governance, institutional accountability and the governance of knowledge. Its work examines how legal, educational and technological institutions construct and exercise evaluative authority, particularly where recognition systems determine access to opportunity, legitimacy and innovation.

The practice publishes interdisciplinary research spanning law, higher education, artificial intelligence and institutional design, prepares submissions to governments and regulators, and develops frameworks for answerability where evaluative decisions increasingly depend upon complex organisational or computational systems. Across these domains, its work seeks to make evaluative authority transparent, reconstructable and institutionally answerable.

Lex et Ratio Ltd is a company registered in England and Wales, Company No 16641457. Registered office: Suite A, 82 James Carter Road, Mildenhall, IP28 7DE, United Kingdom.

www.lexetratio.com

Standards That Can Lose

AN EPISTEMIC-FIDUCIARY CONSTITUTION FOR THE UNIVERSITY

Peter Kahl, Lex et Ratio Ltd
18 July 2026

1. Introduction

1.1 The Problem, the Claim, and What Is New in It

A university recognises some contributions and not others. It admits, appoints, promotes, funds, credits, publishes and declines, and every one of those decisions is made under an evaluative standard: an account, largely unwritten and nowhere assembled, of what counts. The standards are not incidental to the institution. They are the greater part of what makes it a university rather than a set of buildings with a library.

They are also barely examined by those they govern, and — this is the paper’s central contention — only partially examinable by those who hold them, and not exhaustively or independently validatable from within.

Chapters 2 and 3 set out what follows from that in the contemporary university: a drift from epistemic authority toward visual authority, fiduciary duties reduced to their public presentation, and collective epistemic labour resolved into the image of whoever leads. The remainder of the paper asks what could be done about it, and arrives at a question narrower than the one the literature usually asks. Not how a university might see its own standards more clearly — I argue that no institution can do so exhaustively or validate its own standard from within — but **who may compel it to look again, and what happens when it declines**. The answer this paper builds toward is the one its title names: a standard the university may not place beyond reach, a standard that must remain *capable of losing*.

The term, and its prior uses. *Epistocracy* is borrowed and redefined here, not coined, and three usages must be told apart before the argument relies on the word.

The first is not epistocracy at all but its near-namesake, *epistocracy*, the term David Estlund popularised for rule by those regarded as knowing more — ‘rule by the knowers’ rather than rule by the people.¹ Estlund raised it in order to reject it: his point, which is also this paper’s, is that ‘the move from expertise to authority is a fallacy’.² The model proposed here is anti-epistocratic in exactly that sense. Expertise may carry epistemic weight; no expert acquires epistemic finality.

¹ David Estlund, *Democratic Authority: A Philosophical Framework* (Princeton University Press 2008) 3, 30.

² Estlund, *Democratic Authority* 39..

The second is Taleb's. In *The Black Swan* he used 'epistemocracy' for a polity governed 'from the basis of the awareness of ignorance, not knowledge', whose office-holder — the *epistemocrat*, of whom he makes Montaigne the modern exemplar — 'holds his own knowledge to be suspect'.³ The present usage shares Taleb's starting point, that authority cannot secure a sufficient view of its own knowledge, but parts from him on where the safeguard sits. Taleb's epistemocracy is *dispositional*: it asks for epistemocrats, rulers conscious of their own fallibility. This paper expressly refuses to make humility, virtue or carefulness the operative check (§7.3), for the reason the seal supplies — an epistemically humble authority is still the holder of its own standard, and its humility does not reach the residue that reflection is conducted under. The contribution is the conversion of that disposition into a structure:

Taleb's epistemocrat knows the ruler may be wrong. The epistemocracy proposed here assumes that a ruler who knows this *still* cannot identify the full standard through which the error is produced, and therefore gives another party the standing and the institutional means to compel the standard's reopening.

The third usage is the one this paper develops: not rule by those who know, nor rule by those who know that they do not know, but **governance under conditions in which no claim to knowledge carries finality** — the disposition of epistemic humility made structural, and made independent of whether any office-holder happens to possess it. When 'I propose epistemocracy' appears below, it proposes this conception and architecture, not the word.

The claim. Universities should be constituted as epistemic fiduciaries in respect of specified evaluative powers.

The formulation is deliberate and the deliberation is doctrinal. This paper does not assert that universities are fiduciaries of knowledge as a matter of English law. They are not, and the assertion would be easy to defeat.

Fiduciary status is not a description of an institution's importance, or of the trust the public reposes in it. It is a conclusion, reached where a person has undertaken to act for or on behalf of another in a particular matter in circumstances giving rise to a relationship of trust and confidence, and whose distinguishing incident is loyalty: on Millett LJ's classic formulation, 'the distinguishing obligation of a fiduciary is the obligation of loyalty'.⁴ Discretion over another's interests and that other's vulnerability are relevant to the conclusion; they do not compel it. And the order of derivation runs the opposite way to the one an argument like this one is tempted to assume: the obligations are not owed because a person is a fiduciary, but rather the person is a fiduciary because the obligations are owed. Status follows duty. It cannot be used to generate it.

The fiduciary architecture already present in the university therefore points away from this paper's beneficiaries rather than toward them. Where a university is incorporated as a company, its directors' duties are owed to the company.⁵ Where its governing body are charity trustees, their obligations run to the charity's purposes. Neither runs to the applicant it rejected. The relationships this paper is concerned with — with students, applicants, junior colleagues and rejected researchers — are contractual, statutory, regulatory, or nothing at all, and in none of them has the university undertaken to act on the other's behalf in respect of the standard by which it judges them.

³ Nassim Nicholas Taleb, *The Black Swan: The Impact of the Highly Improbable* (Random House 2007) ch 12, 190–92.

⁴ *Bristol and West Building Society v Mothew* [1998] Ch 1, 18 (Millett LJ).

⁵ Companies Act 2006, s 170(1).

That absence is the finding, not the obstacle. The undertaking is what is missing, and supplying it is what this paper proposes. The proposal is that in respect of an identified set of evaluative powers — admission, assessment, credit, promotion, funding, recognition — the university be placed under an obligation to exercise them for their proper other-regarding purposes, with loyalty to the *protected evaluative interest* of the class governed by each power, carrying the one incident that makes fiduciary law worth borrowing here: **an obligation is not discharged by the obligor's own sincere assessment that it has been discharged.**

The protected evaluative interest must be specified, because much of the objection to fiduciary language in this setting comes from misreading it. The interest is *not* an interest in a favourable substantive outcome. The rejected applicant is not owed admission, and the unpromoted colleague is not owed promotion; a fiduciary framing that implied otherwise would be incoherent in every competitive or scarce decision, where to serve one candidate's outcome is to deny another's. What each member of the governed class is owed is narrower and shared without contradiction: an interest in being governed by a standard that is constituted for the proper purpose of the power, applied without concealed institutional self-preference, and maintained in a condition of answerable defeasibility. So construed, the duty generates no zero-sum conflict between beneficiaries, because none of them is owed the position — all of them are owed the same properly-constituted and defeasible standard.

Four questions follow immediately, and the proposal is only as good as its answers.

Who owes it? The corporation, acting through whichever body in fact exercises the evaluative power. The duty attaches to the power and not to the office, which is Miller's powers-based account applied as it is meant to be: the label the institution gives the committee is not the test.⁶

What supplies the undertaking? Not consent, since no rejected applicant is in a position to bargain for one. It must be constituted externally, and the form matters. 'Undertaking', in fiduciary theory, ordinarily connotes a voluntary assumption of responsibility; a statute, by contrast, imposes. The obligation proposed here could be constituted in any of three ways — as a statutory fiduciary duty, as a statutory duty carrying specified fiduciary incidents, or through a condition of registration imposing those incidents — and the paper does not depend on which. This bears on the 'status follows duty' point above: if Parliament imposes answerability obligations, it does not automatically follow that a court will *classify* the resulting relationship as fiduciary. The paper does not need that classification. It needs the institutional incidents — loyalty to the protected interest, and non-self-discharge — however they are constituted. That external constitution is, in any of the three forms, the second reason the mechanisms of Chapter 5 cannot be built by a university acting alone.

Owed to whom, given how heterogeneous they are? Severally, and power by power. A registered student, an unsuccessful applicant, an employed academic and an external researcher stand in materially different legal relationships, and a single undivided duty owed to all of them at once would be unenforceable and probably incoherent. The duty is owed in respect of *each evaluative power* to the class governed by that power: the rejected applicant is owed it in respect of admissions standards and not promotion standards.

How are conflicts between those classes resolved? Mostly they do not arise, because the classes do not share a duty to conflict over. Where they genuinely do — a standard governing two classes at once — the conflict belongs to whoever constituted the standard, which is frequently not the university at all (§5.13).

⁶ Paul B Miller, 'The Identification of Fiduciary Relationships' in Evan J Criddle, Paul B Miller and Robert H Sitkoff (eds), *The Oxford Handbook of Fiduciary Law* (OUP 2019) 366.

What is claimed, then, is a constitution and not a diagnosis of positive law. What fiduciary law contributes is not a cause of action but a design: a body of doctrine that already knows how to govern a party whose own account of its conduct cannot be the measure of that conduct. The object of the proposed trust is not the university's reputation, its rankings, or the public's confidence in it. It is the position of those governed by its standards.

A companion study diagnoses the modern university, at the level of the epistemic constitution, as undergoing systemic *fiduciary breach* — a trustee of public reason turned agent of epistemic capture.⁷ This paper takes that diagnosis as its point of departure and advances a deliberately narrower jurisprudential claim, better suited to a mechanism a court or regulator could operate. It does not assert an existing fiduciary duty owed to the epistemic public at large. It proposes a constituted duty owed to identifiable classes in respect of specified evaluative powers, and it defines the breach correspondingly narrowly: not a failure of candour to the public, but the failure to keep an evaluative standard answerably defeasible by those it governs. Breach, on this narrower construction, is not a matter of institutional bad faith or market capture; it is the condition of an evaluative power exercised without the answerability this paper specifies — a condition that a sincere, well-run, publicly candid institution can be in, and typically is.

What is not new. The mechanisms proposed in Chapter 5 are not novel, and it would be a disservice to the argument to present them as though they were. Standing to compel a review; an adjudicator the reviewed party does not appoint or fund; reasons recorded on refusal; thresholds fixed in advance rather than applied at large; protection against retaliation; review at a date the decision-maker cannot move. This is the ordinary furniture of administrative law, and any public lawyer will recognise it on sight. Judicial review, the tribunal, the ombudsman and the whistleblowing statute are assemblies of exactly these parts, and the university is already subject to several of them.

What is new is correspondingly narrow.

First, the justification. These protections are ordinarily defended as safeguards against error, bias, capture or bad faith — contingent goods, warranted because officials are fallible and occasionally venal, and dispensable in proportion as officials are neither. The argument here is not contingent. Under the sealing result (§1.2), no holder of an evaluative standard can *exhaustively and independently validate* the standard under which their own representing proceeds — however diligently they reflect, a residue remains that reflection conducted under the standard cannot reach. Reopening is then not merely a precaution against the possibility that the decision-maker was careless or biased, defects that better disclosure and reflection can reduce. It is the only route by which the *effects* of that residue can defeat the standard without the residue first having to be completely represented — the check operates through contradiction by later evidence, not through a standpoint outside the standard that no one occupies. That converts due process from prudence into necessity, and it fixes the exact limit of transparency: disclosure and reflection are necessary and do real work, but they cannot be self-validating, and the check must therefore come from a pathway the standard-holder does not control.

Second, the domain. These protections attach, in English law and in university regulation alike, to *decisions*, and hardly at all to *standards*.

The point is not impressionistic. England already operates a compulsory independent review scheme for student complaints, and its jurisdiction is defined by exclusions that map the gap almost exactly. A

⁷ Peter Kahl, 'The Fiduciary Breach of the Modern University: Fiduciary Law and the Epistemic Constitution of the Fourth Estate' (Lex et Ratio Ltd 2025) <<https://doi.org/10.5281/zenodo.17618878>>.

complaint is not a qualifying complaint to the extent that it relates to matters of academic judgment.⁸ It must be brought by a person as a student or former student, so an applicant whose application was rejected or mishandled has no standing at all; nor has a member of staff, employment matters being outside the scheme.⁹ Academic judgment, on the adjudicator's own gloss, is the judgment for which the opinion of an academic expert is essential. It is not identical to the evaluative standard — the paper's own distinction between constituting a standard and applying it (§5.13) forbids collapsing the two — but it is the site at which standards are applied, interpreted, and often locally constituted, and the statutory exclusion therefore covers both the expert conclusion in a given case and, indirectly, the standard exercised through it. The carve-out thus reaches precisely the layer this paper is about, whichever way that layer is described.

So the single independent route that exists is available to one of this paper's four beneficiary classes, and is barred by statute from reaching the one subject matter the paper is about. A student may challenge the application of a marking criterion to their script; no one may compel a university to submit the criterion itself to what has since become of those it excluded. An unsuccessful applicant may allege that a procedure was not followed; none may require the procedure to answer for the fifteen years of people it did not admit. The vacancy is not in the law of decisions. It is one layer up, where the standards are, and that is the layer this paper addresses. The condition is a general one, not peculiar to universities: decision-making authority persists and is exercised while the evaluative framework governing it is operable by no one in the chain — a structure of *delegated discretion* that recurs across fiduciary and administrative law wherever authority is distributed but evaluative control is not.¹⁰

That statutory exclusion is also, incidentally, the best available evidence that the duty proposed here is **new**. Parliament did not omit academic judgment from the scheme by oversight. It carved it out.

Third, the beneficiary. The duty is owed to an identifiable present party — the student, the applicant, the junior colleague, the rejected researcher — and not to truth, to the community of inquirers, or to the next generation (§1.2). This matters practically rather than decoratively: a duty owed to an abstraction is a duty on which no one can sue.

A claim not made. Nothing here asserts that a university adopting these arrangements will be better regarded, or that its adoption would restore public confidence in the sector. Chapter 5.12 argues the opposite, and treats the temptation to argue otherwise as the specific way this proposal will be destroyed.

The structural point that carries the paper, and which the fiduciary framing makes available, is this: fiduciary law already knows how to govern a party who cannot be relied upon to be the final judge of its own conduct. The university is such a party — not because it is corrupt, but because nobody can make their own assessment the final measure of the evaluative standards under which that assessment proceeds, and the university is nobody in particular.

⁸ Higher Education Act 2004, s 12(2).

⁹ Office of the Independent Adjudicator for Higher Education, *Guidance on the Rules*, rr 4–5 <<https://www.oiahe.org.uk/about-us/our-scheme/our-rules/guidance-on-the-rules/rule-4-and-rule-5/>> accessed 17 July 2026.

¹⁰ Peter Kahl, 'Delegated Discretion: When Authority Lacks Evaluative Control in Fiduciary and Administrative Law' (version 2.0, Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20920750>>.

1.2 Three Imported Results

The remedy proposed in Chapters 4 and 5 depends on three results developed elsewhere in this series. They are stated here and not re-argued; the arguments are in the cited works.

The seal. Theoretical representation is conducted within a structured system of admissible distinction — a *law* — and that structure constitutes what such representation can represent. It follows that a domain cannot exhaustively and independently represent the condition that generates its own terms of representation.¹¹

The strength of the claim matters, and it is easy to overstate, so it should be fixed precisely. The seal does *not* assert that an institution can represent none of its evaluative criteria. It plainly can represent some. A promotions committee can come to see that it has been rewarding grant income; an admissions panel can discover that it has treated institutional pedigree as a proxy for ability; a discipline can identify, in retrospect, an entrenched methodological preference. Such reflection is real, and it is valuable, and nothing in this paper disparages it. What the seal denies is that such reflection can ever be *complete* or *self-validating* — that an authority can obtain a standpoint outside its own standard from which the standard's full contribution to its judgements becomes visible and can be certified correct. Every act of reflection is itself conducted under a standard, including the standard governing what counts as a criterion worth noticing; the regress does not terminate in a standard-free vantage. Reflection can reach any given criterion. It cannot reach all of them at once, from nowhere, and confirm the set.

Applied to institutional evaluation, this cuts in both directions. Those governed by a standard cannot exhaustively represent it: they judge under it. And the holders of the standard cannot exhaustively validate it either: it is the law under which their own representing proceeds. The residue that no reflection reaches is not a small remainder. It is likely to contain precisely those operative assumptions most resistant to internal challenge, and may therefore account for otherwise unexplained patterns of exclusion — as the history of disciplines that have revised standards they were once certain of repeatedly suggests. The claim is not that all exclusion issues from the residue; some is generated by explicit and consciously defended criteria, and is contestable in the ordinary way. It is that the exclusion hardest to see, and so most in need of an external route, is disproportionately the residue's.

Nobody has the whole view. No bad faith is required for a hollowed institution to be sincere about its own openness, and — importantly — no obtuseness is required either. An institution can reflect diligently, disclose fully, revise what it finds, and still not have reached the operative residue, because reaching it would require the standpoint the seal denies. This is why the diagnosis of Chapter 3 does not depend on attributing cynicism to anyone, and why the remedy of Chapters 4 and 5 cannot depend on cultivating diligence: the defect survives both.

Transparency is necessary but not sufficient, and not self-validating. Disclosure, explanation, reason-giving and reporting do real work: they constrain arbitrariness, expose the criteria an institution *can* see, make challenge possible, and fix a claim in time so that later evidence can contradict it — §5.5 depends on precisely these functions. What they cannot do is supply the missing check, because the check would require the institution to represent and certify the standard under which its own representing proceeds, and that is what the seal forbids. Transparency renders an institution's account of itself; it cannot render the part of the standard the account is produced under. What remains, beyond representation, is **exposure**: the standard must remain capable of *losing* — defeatable by what happens,

¹¹ Peter Kahl, 'The Inside of Generatedness: Representational Sealing and the Materials of Temporal Order' (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20353457>>.

through a route its holder does not control. This is the content of answerability, and it is why the corpus says *reconstructable* rather than explainable, and *reopenable* rather than merely transparent.

The operative conditions. Three things must be held apart¹²:

- **Procedural authority** — over how things are vetted and when — may and should be concentrated in identifiable hands. Diffusing it produces not accountability but the absence of anyone to hold.
- **Epistemic weight** — the authority that well-grounded judgement carries — is properly distributed, and its distribution is a genuine good.
- **Epistemic finality** — the standing to place a conclusion beyond reopening *because of who holds it* — must exist nowhere.

Where the forms of scrutiny are retained around conclusions the forms may not reach, the result is **hollowed authorisation**. The governing property is single: *the capacity to reopen a conclusion cannot be controlled by the authority whose conclusion would be reopened*. Its minimum conditions are six: independence; standing to trigger; recorded reasons for refusal; protection against retaliation; demanding but non-discretionary thresholds; and periodic review against later evidence.

The beneficiary. The party to whom the fiduciary duty is owed is whoever is currently governed by the standard, bears its costs, and cannot monitor or replicate the deliberation that produced it.¹³ In the university: students, junior faculty, applicants, rejected researchers. Not ‘the community of inquirers’. Not truth. Not the next inquirer. The sealed party is present, identifiable, and can be named — and this matters, because a duty owed to an abstraction is a duty no one can sue on.

1.3 Structure of the Paper

Chapter 2 diagnoses the crisis: optocracy and visual authority, epistemic injustice and exclusion, and the rationalisations the sector offers for both. Chapter 3 names the three pathologies, and states at §3.6 what naming cannot do — a limit that bears on how the vocabulary should be expected to travel.

Chapter 4 sets out the three principles of epistocracy and, at §4.6, ranks them: two govern what an institution can see, and only the third bears on what happens when it cannot. Chapter 5 is the remedy. It opens with the test every mechanism must pass (§5.1), places periodic review against later evidence first (§5.2), sets the four conditions on oversight bodies (§§5.3–5.6), answers the regress objection (§5.7), tests the remaining mechanisms of the standard design against it (§§5.8–5.10), fixes the order of implementation (§5.11), specifies what follows when a standard loses (§5.12) and where the finding must travel when the standard was constituted elsewhere (§5.13), and closes on the defect most likely to destroy the whole proposal (§5.14).

Chapter 6 anticipates resistance, and at §6.5 states the conditions under which this paper is defeated. Chapter 7 situates the argument against agonistic pluralism, the pedagogy of risk, virtue epistemology and Barnett’s ecological university. Chapter 8 asks what transfers beyond the university and on what conditions. Chapter 9 concludes.

¹² Peter Kahl, ‘Answerable Knowledge: Procedural and Epistemic Authority in the Research Organisation’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20439758>>.

¹³ Peter Kahl, ‘Epistemic Humility as Fiduciary Obligation: Entrusted Discretion and Responsibility for Belief’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.18440605>>.

2. Diagnosing the Epistemic Crisis

2.1 Institutional Optocracy and Visual Authority

Contemporary universities increasingly conflate institutional legitimacy with visual prominence rather than substantive epistemic contribution. Senior leaders' images dominate branding campaigns, league-table announcements and ceremonial events, producing what I have termed **optocratic drift**: the gradual substitution of epistemic credibility by institutional optics. This is not a marketing tactic. It is a structural redefinition of legitimacy, privileging the visible over the epistemic.

Foucault's analysis of disciplinary power clarifies the mechanism: visibility is never neutral, but is one of the means by which norms are imposed and knowledge hierarchies enforced.¹⁴ In universities, optics translate into authority. Leadership figures and rankings come to symbolise epistemic credibility, while diffuse and less legible practices — mentoring, collaborative teaching, community engagement, the labour of technicians and librarians — are marginalised, not by decision but by illegibility.

Barnett adds a further dimension. Universities inhabit a condition of **supercomplexity**, characterised by competing and incommensurable frameworks of meaning, and their normative role is to equip individuals and societies to live with, and respond critically to, that condition.¹⁵ Under optocratic drift, institutions retreat from supercomplexity. Plurality flattens into simplified visual narrative; uncertainty resolves into glossy images of 'excellence'. Rather than building resilience in the face of complexity, universities reduce knowledge to surfaces — and thereby fail at precisely the task Barnett identifies as theirs.

The diagnosis to this point is drawn in the register of critical theory, and its empirical warrant lies elsewhere, in a body of institutional research this paper relies upon rather than reproduces. That rankings do not merely observe universities but change them is the documented finding of Espeland and Sauder's study of law-school rankings, which identifies commensuration — the compression of plural qualities onto a single scale — as a mechanism by which a public measure recreates the world it purports to describe.¹⁶ That the metricisation of research assessment carries known costs — to interdisciplinarity, to work that resists the indicator, to the standing of the expert judgement it was meant to support — is the conclusion of the United Kingdom's own independent review of research metrics.¹⁷ And that audit, once installed, tends to reconstitute in its own image the practices it measures is the burden of the anthropological literature on audit cultures.¹⁸ This paper takes these findings as established and asks the question they raise but do not answer: what would make such a standard answerable to those it misrecognises.

A distinction should be marked here that the rest of the paper depends on. Optocratic drift is not a failure of sincerity. The institution producing the brochure believes in it. This will matter greatly in Chapters 4

¹⁴ Michel Foucault, *Discipline and Punish: The Birth of the Prison* (Alan Sheridan tr, Pantheon 1977).

¹⁵ Ronald Barnett, *Realizing the University in an Age of Supercomplexity* (SRHE/Open University Press 2000).

¹⁶ Wendy Nelson Espeland and Michael Sauder, 'Rankings and Reactivity: How Public Measures Recreate Social Worlds' (2007) 113 *American Journal of Sociology* 1.

¹⁷ James Wilsdon and others, *The Metric Tide: Report of the Independent Review of the Role of Metrics in Research Assessment and Management* (HEFCE 2015).

¹⁸ Marilyn Strathern (ed), *Audit Cultures: Anthropological Studies in Accountability, Ethics and the Academy* (Routledge 2000).

and 5, where the temptation is to prescribe more honesty and find that honesty is not the missing ingredient.

2.2 Epistemic Injustice and Exclusion

Reliance on optics entrenches structural epistemic injustice. Fricker’s account of **testimonial injustice** illuminates how credibility becomes unfairly distributed: those with less visual prominence — early-career scholars, community partners, students, precariously employed staff — are systematically disadvantaged.¹⁹ Their contributions, though epistemically significant, carry reduced weight because they are less easily packaged or displayed.

Dotson’s notion of **contributory injustice** sharpens the critique: entire categories of knowledge are excluded when they fail to conform to dominant institutional norms, and the exclusion operates through the wilful hermeneutical ignorance of those in a position of epistemic authority rather than through any explicit refusal.²⁰ Indigenous methodologies, critical pedagogies and decolonial epistemologies frequently suffer this fate. Their resistance to easy representation makes them vulnerable to erasure under optocratic regimes — and, in Dotson’s stronger formulation, epistemic oppression can be *third-order*: irreducible, structural, and invisible from within the resources of the dominant framework.²¹

That last point is worth dwelling on, because it anticipates §1.2 from within the epistemic injustice literature and independently of it. Dotson’s third-order oppression is a case in which the framework’s own resources cannot register what the framework excludes. This is the seal, arrived at from a different direction, and it is the reason the remedy in this paper cannot take the form of asking the institution to look harder.

Medina’s work on **resistant epistemologies** provides a counter-narrative: marginalised knowledge practices can and must be cultivated as forms of epistemic resistance.²² Yet under optocratic drift the opportunities for resistance shrink. The result is a systematic narrowing of the epistemic horizon, in which plurality is tolerated exactly insofar as it can be visualised, branded, or monetised.

There is a further mechanism, and it bears directly on the remedy. Metric governance makes recognition scarce, and scarcity operates vertically as well as horizontally: it produces **epistemic clientelism**, the moderation of independent judgement so as to remain legible to the patrons who control advancement.²³ The clientelist does not withhold an objection because anyone has silenced them. They withhold it because voicing it is illegible to the standard by which they will be advanced — and they can see *that* much, even though they cannot see the standard itself.

Note what this does to any remedy built on voice. The parties best placed to notice that a standard excludes are the parties with the least room to say so. Their silence is not evidence of consent, and an institution that reads it as consent is not being dishonest. It is reading the equilibrium correctly and drawing the wrong conclusion from it.

¹⁹ Miranda Fricker, *Epistemic Injustice: Power and the Ethics of Knowing* (OUP 2007).

²⁰ Kristie Dotson, ‘A Cautionary Tale: On Limiting Epistemic Oppression’ (2012) 33 *Frontiers: A Journal of Women Studies* 24.

²¹ Kristie Dotson, ‘Conceptualizing Epistemic Oppression’ (2014) 28 *Social Epistemology* 115.

²² José Medina, *The Epistemology of Resistance: Gender and Racial Oppression, Epistemic Injustice, and Resistant Imaginations* (OUP 2013).

²³ Peter Kahl, ‘The Recognition Game: Metric Governance and the Stabilisation of Epistemic Clientelism’ (version 2.0, Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20774711>>.

Biesta underscores the educational cost. Genuine education requires risk, unpredictability and openness to the unknown²⁴, and these are systematically suppressed when institutions prioritise legibility and control. Mouffe's theory of agonistic pluralism highlights how disagreement and contestation are vital to democratic vitality.²⁵ Optocratic universities sanitise conflict into consensus imagery, and thereby undermine the pluralism they claim to sustain.

2.3 Institutional Rationalisations

Universities justify their reliance on visual practice in terms of transparency, accountability and market competitiveness. Public visibility, the argument runs, demonstrates leadership responsibility, facilitates trust, and strengthens fundraising. This is the logic Slaughter and Rhoades describe as **academic capitalism**: the fusion of higher education with market-oriented practice and competitive branding.²⁶

These rationalisations mask deeper exclusions. Visibility of leaders may appear to enhance transparency, but it selects for what displays; fundraising campaigns and league-table optics may attract resources, but they do so by narrowing the epistemic mission to marketable outputs.

There is a second and less obvious problem with the rationalisation, and it is the one this paper pursues. *Transparency is the wrong good*. The rationalisation and the standard critique of it share a premise: that what the institution owes is a faithful account of itself, and the dispute is only over whose face appears in it. Chapter 4 rejects the premise. An institution can give a faithful account of itself, in complete good faith, and the account will be silent about exactly the thing at issue — because the standard under which it selects what to disclose is the thing at issue, and it is disclosing under that standard.

What the university owes is not a better picture. It is exposure: standards that can lose.

3. Naming the Pathologies

To resist the crisis outlined in Chapter 2, universities require a sharper vocabulary, and this chapter supplies three terms: optocratic drift, fiducial hollowing, and epistemic inversion.

The claim made for them should be stated at the outset, because it is narrower than the claim usually made for critical vocabulary. **Naming does not enable an institution to see what it is doing**. Under the seal (§1.2) it cannot: an institution whose evaluative standard makes its optics look like communication is not assisted by being handed the word for that condition, and will reach for the word to describe its competitors.

What naming does is smaller, and more useful. It gives the party a standard excludes a way to report a pattern without having to represent the standard — and a pattern is what the mechanisms of Chapter 5 run on. The vocabulary is an instrument of the sealed party, not a mirror for the institution. §3.4 states the duty on that footing, and §3.6 develops what follows from it.

²⁴ Gert Biesta, *The Beautiful Risk of Education* (Routledge 2013).

²⁵ Chantal Mouffe, *Agonistics: Thinking the World Politically* (Verso 2013).

²⁶ Sheila Slaughter and Gary Rhoades, *Academic Capitalism and the New Economy: Markets, State, and Higher Education* (Johns Hopkins University Press 2004).

3.1 Optocratic Drift

Optocratic drift is the incremental prioritisation of institutional visibility over substantive epistemic integrity. Resources flow to league-table positioning, glossy branding and leader-centric publicity rather than to rigorous and diverse scholarly practice. The positioning is worth having because commensuration has made it worth having: once academic standing is projected onto a single scale, recognition becomes scarce, and a scarce positional good will be competed for whatever anyone believes about its validity.²⁷

Over time, drift normalises the association of credibility with optics: to be seen is to be believed. This marginalises epistemic practices that resist display — mentoring, collaborative teaching, community-engaged scholarship, the slow correction of a literature. The word ‘drift’ is chosen deliberately: no one decides on optocracy. It accumulates from decisions each of which is defensible, and the absence of a decision is what makes it hard to reverse and impossible to confess.

The mechanism has a precise form once a metric enters. When a university optimises for a research-assessment metric, it does not, past a certain point, work *through* the metric toward the good the metric was meant to track. It works *for the metric itself*: the proxy becomes the object served, and the underlying good becomes whatever the proxy happens to capture. This is more than Goodhart’s observation that a measure under pressure ceases to measure well, and more even than the reactivity Espeland and Sauder describe, in which the measured change their conduct under observation²⁸; it is a transfer of the institution’s *ends*. The metric’s constitutor — the assessment body, the ranking, the funder — has thereby acquired the university’s evaluative allegiance without acquiring any answerability for how the university treats those its pursuit of the metric excludes. This is *capture of the university’s evaluative allegiance by the metric’s constitutor*: the specific, tractable form taken, at the level of a single evaluative power, by the epistemic capture a companion study diagnoses at the level of the institution as a whole.²⁹ That the allegiance runs to the metric and not through it is why, when a standard so shaped is defeated, the finding must travel to the metric’s constitutor and not stop at the university that served it (§5.13).

3.2 Fiducial Hollowing

Fiducial hollowing is the erosion of fiduciary duties into symbolic gesture.

Fiduciary law, as Frankel established, demands care, loyalty, transparency and accountability.³⁰ Smith develops the point furthest: fiduciary principles govern the use of legal powers held for other-regarding ends, in private law and public law alike, and their core requirement is that such powers be used only for the purposes for which they were granted, and never where the holder stands in a conflict of interest and duty.³¹ Miller shows that fiduciary relationships are identified by the presence of discretionary power over another’s practical interests, however the relationship is labelled by those inside it.³²

²⁷ Kahl, ‘The Recognition Game’.

²⁸ Espeland and Sauder (n above) — commensuration and self-fulfilling prophecy as the mechanisms of reactivity.

²⁹ Peter Kahl, ‘The Fiduciary Breach of the Modern University: Fiduciary Law and the Epistemic Constitution of the Fourth Estate’ (Lex et Ratio Ltd 2025) <<https://doi.org/10.5281/zenodo.17618878>>.

³⁰ Frankel, *Fiduciary Law*.

³¹ Smith, *The Law of Loyalty*.

³² Miller, ‘The Fiduciary Relationship’; Miller, ‘The Identification of Fiduciary Relationships’.

Keech v Sandford exemplifies the strictness.³³ A trustee took the renewal of a lease that the trust could not itself obtain. No harm to the beneficiary was shown; none could be. The trustee was nonetheless required to hold the lease on trust and account for the profit. Lord King's reason is the one that matters here: it may seem hard that the trustee is the only person of all mankind who might not have the lease, but it is *very proper* that the rule be pursued and not in the least relaxed. The duty of loyalty is strict, and no accounting of good consequences licenses its subordination.

The rule has not softened. In 2025 the Supreme Court declined an invitation to admit a 'but for' counterfactual into the fiduciary's duty to account — the defence that the profit would have been made anyway, without the breach — holding that the rule is prophylactic and that treating it as a mere remedy would water down its effect.³⁴ That refusal is worth pausing on, because the counterfactual equity declines to entertain is one constructed by the party who benefits from it. §5.2 declines an isomorphic counterfactual for an isomorphic reason.

The rule is strict because the fiduciary cannot be relied upon to assess their own conduct — not because fiduciaries are dishonest, but because the assessment is made under the interest it is assessing.

That is an analogy to the seal and not a derivation from it, and the distinction should be marked rather than blurred. *Keech* and *Rukhadze* are about conflicts of interest; the seal is about conditions of representation, and equity was legislating for neither seals nor standards. But the structural move is the same, and it is old: where a party's own account of a matter is produced under the very thing at issue, the law does not improve the account. It removes the party's control over the outcome.

In higher education, these obligations are routinely reduced to public relations. Universities issue diversity statements, compliance reports and accreditation claims designed for external optics rather than substantive accountability.³⁵ When institutions substitute stewardship with performative compliance, they reproduce what fiduciary law condemns as a breach of loyalty: the subordination of the beneficiary's interest — here, the position of those governed by the institution's evaluative standards — to the institution's self-promotion.

Fiducial hollowing is thus more than oversight. It is a breach of stewardship. And its most important property, the one on which the whole of Chapter 5 turns, is that **it does not require anyone to intend it**.

3.3 Epistemic Inversion

Epistemic inversion occurs when the internal plurality of a university is collapsed into a singular external figurehead. Epistemic labour involving dozens or hundreds of contributors is distilled into the image of a Vice-Chancellor or President; achievements are represented as the work of individuals, and the collective fades from view.

Barnett argues that universities should help societies live with supercomplexity.³⁶ Epistemic inversion undermines that mission by reducing complexity to caricature and plurality to personality. It reinforces a hierarchical optics of knowledge in which visibility substitutes for legitimacy.

³³ *Keech v Sandford* (1726) Sel Cas Ch 61.

³⁴ *Rukhadze v Recovery Partners GP Ltd* [2025] UKSC 10.

³⁵ Peter Kahl, 'Epistemic Justice and Institutional Responsibility in Academia: Toward a Comprehensive Framework for Epistemic Justice in Higher Education' (2nd edn, Lex et Ratio Ltd 2025) <<https://doi.org/10.13140/RG.2.2.11063.38563>>.

³⁶ Barnett, *Realizing the University*.

3.4 Naming as Epistemic Duty, and to Whom It Is Owed

Naming is normative, not merely descriptive. Medina stresses that naming injustice constitutes epistemic resistance, equipping those excluded with the conceptual tools to articulate the exclusion and demand reform³⁷; Zagzebski identifies the intellectual virtue that public naming requires, which is courage.³⁸

Both accounts are addressed to a knower who can, with effort, come to see the injustice. The present argument cannot borrow them at that width. Under the seal, the exclusion is not something the excluded can articulate either: they cannot state the standard that excluded them, the panel that applied it cannot state it, and no quantity of courage supplies what neither party can see. A vocabulary that asked the excluded to name the standard would be asking them for the one thing their position makes unavailable.

What survives the narrowing is precise, and is enough. Medina's excluded knower does not need the concept in order to perceive that something is wrong; they need it in order for their report of it to be **admissible**. Without a word, a rejected researcher's account of their treatment is anecdote — a complaint about an outcome, which any institution can absorb, and which clientelism (§2.2) makes costly to voice at all. With a word, the same account states a class: not *I was treated unfairly*, which invites a reply about the merits, but *this has the shape of contributory injustice, and the recognition system has registered no work of this kind in nine cycles*, which does not. The second is a pattern claim. Pattern claims meet the threshold at §5.3; complaints do not.

Naming is therefore a fiduciary duty of the institution, rather than an epistemic virtue of the individual, and it is owed for a reason that has nothing to do with the institution's own self-understanding. **The institution owes its beneficiaries the vocabulary in which they can compel it to answer.** That the vocabulary will not help the institution see is no objection to owing it: fiduciary duties are owed to beneficiaries, not to fiduciaries.

3.5 Scholarly Continuity

These terms extend and integrate earlier work. In 'Directors' Epistemic Duties and Fiduciary Openness' I argued that fiduciary law requires epistemic responsiveness, not merely structural compliance — a requirement hollowed out in higher education governance.³⁹ In 'Epistemic Justice and Institutional Responsibility in Academia' I set out how institutions suppress uncertainty to preserve reputational control, and thereby commit epistemic injustice while reporting accountability.⁴⁰

In 'The Recognition Game' I traced the mechanism that produces both.⁴¹ Metric governance does not merely measure research; it governs recognition. Commensuration of academic standing onto a single scale makes recognition **scarce**, and scarcity has two effects at once. Horizontally, it adversarialises: colleagues become rivals for a positional good. Vertically, it produces **epistemic clientelism** — the moderation of independent judgement in order to remain legible to the patrons who control advancement.

³⁷ Medina, *The Epistemology of Resistance*.

³⁸ Linda Trinkaus Zagzebski, *Virtues of the Mind: An Inquiry into the Nature of Virtue and the Ethical Foundations of Knowledge* (CUP 1996).

³⁹ Peter Kahl, 'Directors' Epistemic Duties and Fiduciary Openness: A Cross-Cultural and Interdisciplinary Framework for Corporate Governance' (2nd edn, Lex et Ratio Ltd 2025) <<https://doi.org/10.13140/RG.2.2.28155.17449>>.

⁴⁰ Kahl, 'Epistemic Justice and Institutional Responsibility'.

⁴¹ Kahl, 'The Recognition Game'.

That result is imported here and used at §§2.2, §5.2, §5.11, §5.14, §6.2 and 6.3. It earns its place for a reason beyond continuity: it shows that the degraded order is an **equilibrium**. No one need act in bad faith; few would choose the culture that results; and the better culture is held out of reach by the risk of pursuing it alone. This is a second route to the position of §1.2, and it does not travel through the seal. The first says the institution cannot see. The second says that an institution which saw could not, by itself, move.

3.6 The Limits of Naming

The limit was stated at the head of this chapter. It is developed here because it governs how this vocabulary should be expected to travel, and because a paper that levels the audit objection at everyone else owes an account of what its own conceptual apparatus cannot do.

Naming is a demand for representation. It supplies a word for a thing so that the thing may be pointed at. But the pathologies named here are not, in the relevant sense, *pointable* by those inside them. An institution in optocratic drift has, by hypothesis, an evaluative standard under which its optics look like communication and its exclusions look like standards. Handing that institution the word 'optocratic drift' does not give it the ability to see the drift. It gives it the ability to say the word.

What it will do with the word is not a matter for speculation. Every governance vocabulary of the last forty years — transparency, accountability, stakeholder engagement, inclusion, responsible metrics — entered the sector as a critique and left it as a heading. The mechanism is not hypocrisy, and it is not even carelessness. It is §1.2 once more: the institution adopts the word *under the standard the word was coined to indict*, and the adoption is sincere. Sincerity is what makes it useless.

The obvious recommendation therefore has to be refused. Nothing here asks that these terms be written into mission statements, strategic plans or values frameworks in order to clarify institutional self-understanding — the recommendation that critical vocabularies ordinarily carry, and the one an institution encountering this chapter will be readiest to accept. A mission statement is the optocratic reception of an anti-optocratic critique: the institution has taken a diagnosis of its preference for the legible and made it legible. This paper asks nothing of institutional vocabulary at all. It asks that the vocabulary be available to the people the institution's standards exclude, and that a mechanism exist in which they can use it (§5.3).

So: naming lowers the cost of triggering a mechanism. It is not the mechanism. Where a governance vocabulary is adopted and no mechanism accompanies it, the adoption is not partial progress — it is the pathology, operating on the critique.

I therefore record a prediction, so that this paper can lose: of everything proposed here, the neologisms will be adopted most readily and the mechanisms of Chapter 5 least readily, and institutions will use the former to describe their competitors. The prediction recurs in its sharpest form at §5.9, and is listed among this paper's defeat conditions at §6.5.

4. Proposal: Toward an Epistemocracy

4.1 From Pathology to Remedy

The preceding chapters diagnosed the crisis of higher education as marked by optocratic drift, fiducial hollowing and epistemic inversion: epistemic integrity displaced by visual authority, fiduciary responsibility hollowed into symbolic gesture, pluralism collapsed into hierarchical figurehead. Having named these failures, I turn to remedy.

I propose **epistemocracy** — the term borrowed and redefined as set out at §1.1 — a reorientation of higher education governance integrating fiduciary ethics with democratic epistemology and virtue epistemology. Epistemocracy is not a marginal reform. It restructures how universities conceive authority, accountability, and epistemic legitimacy — and it does so by giving up one thing that accounts of institutional accountability generally take for granted.

What has to go is the idea that the remedy consists in the institution's improved self-knowledge. The natural proposal — and the one the sector reaches for whenever the diagnosis of Chapters 2 and 3 is accepted — is a body that would look carefully, report honestly, and represent the institution to itself and to the public more faithfully than the brochure does. Under §1.2 that is not enough: the standard's holder cannot validate, from within, the residue of the standard their own representing proceeds under, and a body appointed by the standard's holder inherits the same blind spot. Improved self-knowledge is worth having and is not the check; what supplies the check is exposure — the standard's liability to defeat by what happens, through a route its holder does not control.

Everything in Chapter 5 follows from that substitution, and the mechanisms that do not survive it are marked as failing where they appear.

4.2 Conceptual Foundations

Epistemocracy draws on four traditions, and takes from each of them something narrower than a synthesis usually claims. The narrowing is not modesty. Three of the four furnish the content of good judgement, and one furnishes an obligation that survives the obligor's sincere belief that they have discharged it. Only the fourth bears the remedy, and a reader is entitled to know which is which before the argument begins rather than after.

Democratic epistemology. Anderson argues that democratic governance requires institutions to embrace diverse epistemic perspectives, empowering individuals to participate meaningfully in collective deliberation.⁴² Epistemocracy operationalises this by decentralising epistemic *weight* and embedding plural perspectives in decision-making. It does not follow, though it is routinely inferred, that decentralised deliberation dissolves finality. A democratic body can finalise. Most do.

Critical pedagogy. Freire emphasises the liberatory potential of education through dialogue, collective empowerment and critical consciousness⁴³; Darder extends this to the institutional reforms required to cultivate inclusive pedagogical communities.⁴⁴ Epistemocracy takes from them the identification of the

⁴² Elizabeth Anderson, 'The Epistemology of Democracy' (2006) 3 *Episteme* 8.

⁴³ Paulo Freire, *Pedagogy of the Oppressed* (Myra Bergman Ramos tr, Continuum 1970).

⁴⁴ Antonia Darder, *Reinventing Paulo Freire: A Pedagogy of Love* (2nd edn, Routledge 2017).

beneficiary: the party whose knowledge the institution is structured not to hear. Freire's contribution here is not the pedagogy but the *addressee*.

Fiduciary ethics. Frankel, Smith and Miller develop fiduciary theory around loyalty, care and accountability in relationships of dependence.⁴⁵ Epistemocracy transposes these into higher education, proposing that universities be constituted as fiduciaries in respect of their evaluative powers. Smith's framing is the most directly useful: the powers at issue are held for other-regarding ends, and the question is always whether the power was used for the purpose for which it was granted — a question the power's holder is structurally the wrong party to answer.

Virtue epistemology. Fricker, Zagzebski and Medina emphasise intellectual virtues — humility, justice, courage — as foundational to responsible knowledge practice.⁴⁶ Epistemocracy incorporates these as constituents of good judgement. It does not, and this is the sharpest departure from how these traditions are ordinarily deployed in governance argument, treat them as the check. A virtue is a disposition of a knower; under the seal, the knower cannot see the thing about which the virtue would be exercised.

Courage will not show a standard's holder what the standard excludes. Humility is not a substitute for a mechanism, and an institution that is humble about its blind spot has not thereby seen it.

The three principles. From this synthesis epistemocracy derives three principles — distributed credibility, opacity-respecting pluralism, and fiduciary answerability. They are not co-equal, and §4.6 says why.

Synthesis, and its division of labour. Democratic epistemology grounds plurality; critical pedagogy identifies the beneficiary; virtue epistemology furnishes the qualities of good judgement. Each is necessary to an institution worth having, and none of them is a check. A plural body can finalise, a beneficiary can be identified and still hold no key, and a virtuous office-holder cannot see the standard under which their virtue is exercised (§7.3). Fiduciary ethics supplies what the other three cannot: an obligation that survives the obligor's sincere belief that they have discharged it, and a beneficiary entitled to enforce it.

This division of labour should be held in view through Chapters 4 and 5, because it explains the shape of what follows. The first three traditions supply the paper's *diagnosis* and its account of what a university is for; they are why the argument is addressed to students, junior colleagues and rejected researchers rather than to abstractions, and why it treats plurality and opacity as goods rather than as inconveniences. They do not supply a single mechanism. Every mechanism in Chapter 5 descends from fiduciary law and the seal, and a reader who finds Freire, Mouffe and Zagzebski doing less work in the second half of this paper than the first has read it correctly. **The traditions tell us what is owed and to whom. Only one of them tells us what happens when the debtor cannot see the debt.**

4.3 Distributed Credibility

Distributed credibility requires universities to recognise epistemic contribution across the full range of institutional actors. Under optocratic drift, recognition accrues disproportionately to senior administrators and highly visible scholars, and metric commensuration compounds the concentration by

⁴⁵ Frankel, *Fiduciary Law*; Smith, *The Law of Loyalty*; Miller, 'The Fiduciary Relationship'; Miller, 'The Identification of Fiduciary Relationships'.

⁴⁶ Fricker, *Epistemic Injustice*; Zagzebski, *Virtues of the Mind*; Medina, *The Epistemology of Resistance*.

making recognition scarce: what is distributed unevenly is also, by construction, in short supply.⁴⁷ Distributed credibility shifts value away from narrow, hierarchical, visibility-based metrics toward the broader ecology of epistemic labour.

The principle aligns with Freire’s pedagogy of empowerment and Darder’s call for structures that validate marginalised voices⁴⁸, and extends Anderson’s democratic epistemology⁴⁹ by ensuring that all contributors — students, adjuncts, technicians, community partners — are credited for sustaining knowledge production. Practical mechanisms include diversified promotion criteria, recognition of mentoring and community engagement, and credit-sharing protocols for collaborative work.

This is a principle of **epistemic weight**. It improves what the institution can see. It is not a check, and §5.8 states what it can and cannot carry.

4.4 Opacity-Respecting Pluralism

Opacity-respecting pluralism affirms the legitimacy of epistemic practices that resist reduction to rankings, metrics or visibility. Many forms of knowledge — indigenous epistemologies, artistic practice, embodied traditions, the tacit knowledge of a technician — cannot be quantified without distortion. Epistocracy preserves their dignity by recognising their validity without forcing assimilation into reductive schemes.

The principle is indebted to Medina’s resistant epistemologies.⁵⁰ Operationalising it involves institutional support for methodologies and outputs that do not conform to conventional measures: narrative reporting, qualitative assessment, recognition of non-traditional scholarly outputs.

There is a connection to §1.2 that should be made explicit, because it converts this principle from a gesture of tolerance into a structural claim. Opacity is not a regrettable feature of certain exotic knowledge practices. It is the general condition. *Every* evaluative standard is opaque to its holder; the practices named here are simply the ones whose opacity is visible from outside. An institution that grants opacity-respect to indigenous methodology while assuming its own promotion criteria are transparent has not understood the principle. It has patronised it.

This too is a principle of epistemic weight.

4.5 Fiduciary Answerability

Transparency is a demand for representation: disclose the standard, explain the decision, make the criteria explicit, report against them. It does real and necessary work, and this paper does not disparage it. But it cannot be the whole of the duty, and the reason is structural rather than a matter of institutions withholding. An institution may disclose everything it can see, in complete good faith — and what it can see is, by the seal, not the whole of the standard under which its own seeing proceeds (§1.2). The disclosure will be accurate as far as it reaches. The residue it does not reach is the part that was doing the excluding.

⁴⁷ Kahl, ‘The Recognition Game’.

⁴⁸ Freire, *Pedagogy of the Oppressed*; Darder, *Reinventing Paulo Freire*.

⁴⁹ Anderson, ‘The Epistemology of Democracy’.

⁵⁰ Medina, *The Epistemology of Resistance*.

This is the most uncomfortable consequence of the diagnosis, and the one it is hardest to hold on to: **fiducial hollowing does not require a hollow institution**. A duty discharged as transparency alone is hollowed not by insincerity but by construction, because the operative residue lies outside what any amount of sincere disclosure can render. The reports will be accurate. The reasons will be given. The residue will remain unreached — which is why disclosure must be *supplemented*, not replaced, by exposure.

What replaces it is not more disclosure but exposure. **Fiduciary answerability** is the duty to keep one's evaluative standards capable of losing: defeasible by what happens to those governed by them, through a route one does not control. It is discharged not by an account of the standard but by two things — the standing of the sealed party to compel the standard's reopening, and the standard's periodic submission to what has since occurred.

Two consequences follow, and between them they carry the rest of the argument.

First, disclosure retains a function, but not the one usually assigned to it. Recorded reasons are **evidential**, not explanatory (§5.5).

Second, accountability mechanisms are no longer to be assessed by what they reveal, but by what they can defeat.

4.6 How the Three Principles Relate

Each principle answers a distinct pathology:

- **Distributed credibility** resists epistemic inversion by broadening recognition — a principle of epistemic weight.
- **Opacity-respecting pluralism** counters optocratic drift by legitimising non-visualised practice — a principle of epistemic weight.
- **Fiduciary answerability** remedies fiducial hollowing by keeping standards defeasible — the only principle of answerability among the three.

It is tempting to present the three as interlocking and co-equal, and the temptation should be resisted. They are not co-equal. The first two improve what an institution can see. The third is what operates when it cannot see, which is always. An epistemocracy that adopts the first two and not the third has adopted a better vocabulary for its existing practice, and §3.6 predicts that this is exactly what will be adopted.

5. Institutional Mechanisms

5.1 The Test Every Mechanism Must Pass

The mechanisms below are stated so as to be testable, and the test is the same in each case:

Does this make the standard capable of losing, or does it make someone careful?

Only the first is load-bearing. Any mechanism whose operation depends on an office-holder being vigilant, open, humble or attentive has already failed — not because office-holders are ill-motivated, but because under the seal they cannot see the thing about which they would be vigilant. **One cannot audit one's own blind spot.**

The test disposes of a good deal that is ordinarily proposed under the heading of accountability, and the disposal is the point. The reader should apply it to every paragraph that follows, including the ones I have argued for.

5.2 Periodic Review Against Later Evidence

This mechanism is placed first because it is the only one that does not require anyone to represent what the seal forbids them to represent.

Every other route runs through a challenger who must formulate an objection: state what the standard was, and why it was wrong. The narrowed seal (§1.2) does not make that impossible in every case — a challenger may well articulate a criterion the institution *can* also see, and where they can, the ordinary mechanisms of §§5.3–5.6 do their work. But three things make the challenger-driven route systematically insufficient as the *primary* one, and each is independent of the others.

First, the residue. The part of a standard that reflection cannot reach is, by construction, the part a challenger cannot articulate either — and it is *likely* to be where the more intractable exclusion sits, since a criterion visible enough to be named is a criterion visible enough to have been contested already. Explicit criteria produce exclusion too, and that exclusion is challengeable in the ordinary way; what the challenger-driven route systematically misses is the exclusion generated below the level anyone can name. The objections that can be formulated are drawn from the represented layer; the exclusion that most needs an external route concentrates in the unrepresented one.

Second, clientelism. Even where an objection *could* be formulated, those best placed to formulate it — the junior, the precarious, the applicant hoping to reapply — are those whose advancement depends on the patrons whose standard they would be impugning.⁵¹ A route that runs on challengers selects against exactly the challengers positioned to see, and does so most efficiently where a standard is most entrenched.

Third, capping. The work most likely to reveal a standard's blindness is the work that standard misrecognises, which by definition its holders and its beneficiaries are least equipped to bring forward as evidence.⁵²

A remedy resting primarily on well-formed objection therefore waits for what is, across these three, systematically unavailable — and it will read the resulting silence as the absence of grievance. This is how a hollowed institution acquires evidence of its own health.

Outcomes are under no such constraint. A cohort's trajectory; a rejected line of research that proves fertile elsewhere; a category of contribution a recognition system fails to see in cycle after cycle — these defeat a standard without anyone having had to say what the standard was. Given the seal,

⁵¹ Kahl, 'The Recognition Game'.

⁵² Kahl, 'The Recognition Game'.

outcome-over-time is the most reliable defeater available, and it is therefore the primary mechanism of epistocracy rather than a supplement to the audit board.

In operation:

- **Review dates are fixed at adoption.** Every recognition standard, communication protocol and curricular decision within the framework carries a stated review date, set when the standard is adopted and not extendable by the standard's holder.
- **The burden runs against the standard — once a class-level pattern is shown.** At review, the question is not whether a challenger has made out a case for change, but whether the accumulated record of outcomes is consistent with the standard's continuation. The burden shifts to the standard at the point a repeated, class-level pattern appears (tier two below); a single instance triggers examination without yet shifting anything.
- **The review must track those the standard excluded.** This condition is the whole of the mechanism, and it is the one that will be quietly dropped in implementation. A review that examines only those admitted, promoted, funded or recognised examines only the standard's successes, and will confirm the standard every time — with data, at length, and sincerely. Rejected applicants, unpromoted staff, unfunded proposals, and contributions the credit system never registered are the evidentiary base. Those admitted are the control, not the sample.
- **Watch for out-of-distribution failure.** The signal is systematic misperformance on cases lying outside the distribution that formed the standard. This is the same claim in machine-learning clothing: an evaluative system that fails out-of-distribution is not encountering hard cases but exhibiting the boundary of its own enclosure, and the failure is structural rather than statistical.⁵³ A university whose recognition system performs well on candidates resembling its existing staff and poorly on everyone else has not found that everyone else is worse — but establishing that requires the comparator described below, not merely the observation of poorer outcomes.
- **Expect the failure to present as quality.** A recognition system under metric governance exhibits *evaluative-layer capping*: it rewards work extending inherited standards and misrecognises the generative work that would reopen them.⁵⁴ The misrecognition is not noise in the system; it is the system's competence. So the reviewer should not expect a defeated standard to announce itself as an error. It will present as a body of work the institution declined — correctly, by its own lights, with reasons, at the time — and which went on to matter somewhere else.
- **Findings and the evidentiary base are published**, and where a standard *survives* review the grounds are recorded — because survival is itself a determination, and a determination no later evidence can touch is finality under another name.

Periodic review does not require anyone to be careful. It requires the institution to have committed, in advance, to a date on which it does not control what the evidence will say.

The evidentiary base: feasibility and lawfulness. A mechanism that turns on the later trajectories of people an institution rejected invites two objections at once — that the data does not exist, and that collecting it would be unlawful. Both are serious, and an accountability proposal that requires no

⁵³ Peter Kahl, 'Out-of-Distribution Failure as a Structural Signal of Epistemic Enclosure' (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.19353400>>.

⁵⁴ Kahl, 'The Recognition Game'.

evidence never has to meet them, which is part of why such proposals are preferred. They are met here rather than left for implementation, since a mechanism whose evidentiary base is unspecified is a mechanism that will be piloted without one.

What the review needs, and what it does not. The review asks a question about a standard, not about a person. It needs to know whether a class of contribution has gone unrecognised across a stated number of cycles — not what became of any particular applicant. That distinction is not a nicety. It determines everything that follows, because a body seeking aggregate answers about cohorts needs neither individual-level disclosure to the institution nor any decision about an individual, and a body seeking dossiers needs both and should not have either.

Where the data is. Most of it is already held. Promotion, funding and curricular standards govern populations that remain in view: staff who were not promoted are still employed, proposals that were not funded resurface elsewhere, students subject to a curricular decision are still enrolled. The contemporaneous records required by §5.5 and §5.8 — refusals with grounds, credit as contributed — supply the rest, provided they are kept from the first day rather than reconstructed at review (§5.11). The genuinely hard case is admissions, where the rejected leave and the institution has no further relationship with them.

For that case the sector already operates much of the machinery, though pointed in one direction. The Longitudinal Education Outcomes dataset links education records to tax and benefit records, and — unlike the graduate-outcomes survey, which is derived from student records and targets graduates — it already incorporates UCAS *application* data, so the linkage of applicants to later outcomes is not a fanciful proposition but an extension of a structure that exists. Every established use of it, however, observes those admitted or enrolled. The point to be careful about is that a lawful basis for processing is not a data-access gateway: the research-and-statistics provisions permit processing, but they do not by themselves confer access to UCAS, HESA, HMRC or DWP records. Extending observation to the rejected population would require an express legal basis for access, data-sharing agreements between the controllers, and a secure processing environment — a substantial governance and linkage project, not a minor operational adjustment. What can fairly be said is narrower and still significant: the apparatus exists, the applicant-to-outcome linkage is already technically instantiated in LEO, and the reason it has not been turned on the excluded is not primarily technical or evidential. **The machinery exists, and it is pointed at the standard's successes.**

The lawful basis. Under UK data protection law the difficulty is not insurmountable, and its resolution is instructive about the mechanism's own design. The framework here is the UK GDPR and the Data Protection Act 2018 as amended by the Data (Use and Access) Act 2025, which renumbered the safeguards regime for research, archiving and statistical purposes — the 'RAS purposes' — from Article 89 into new Articles 84A to 84C, and inserted a statutory definition of processing for statistical purposes as processing whose results are produced in aggregate form and are not used in support of decisions about any particular data subject.⁵⁵

Three points follow, and they are cumulative rather than alternative. First, *compatibility and retention*: retaining application data in order to evaluate the admissions standard is further processing for a purpose other than collection, which engages the purpose-limitation principle — but processing for the RAS purposes, subject to the safeguards in Articles 84B–84C, is treated as compatible with the original purpose, and the storage-limitation principle is relieved to the extent longer retention is genuinely

⁵⁵ Data (Use and Access) Act 2025, ss 67, 85, inserting UK GDPR arts 84A–84C and amending art 4; the statistical-purposes condition tracks the aggregate-only definition now in art 4.

needed for those purposes. Second, a *lawful basis under Article 6* remains independently necessary; compatibility is not itself a basis. A university acting as a public authority relies on the public-task basis, legitimate interests being unavailable to it for processing in performance of its tasks.

Third, and this corrects a conflation that would otherwise weaken the argument: an additional *Article 9* condition is required only where the data are genuinely special-category data, and **not every Equality Act protected characteristic is special-category data**. Race, ethnicity, religion or belief, and sexual orientation fall within Article 9; sex, age, and marital status do not, though they are protected characteristics. Where the excluded class is defined by an Article 9 characteristic — or by a proxy from which one can be inferred, which the regulator treats as engaging Article 9 — the research-and-statistics condition in Schedule 1 to the Data Protection Act 2018, read with the RAS safeguards, supplies the additional basis; where the class is defined only by a non-special characteristic, Article 6 alone suffices. The paper should not, and now does not, assume that every excluded class triggers the special-category regime.

Transparency is discharged at the point of application, in the privacy notice, which should say plainly that the standard applied will itself be reviewed and that the applicant's outcome forms part of the evidence.

The instructive part is the condition built into the definition of statistical processing itself: it is available only where the results are produced in aggregate and are *not used to support measures or decisions with respect to a particular data subject*. §5.2 satisfies that condition exactly, and not by accident. It reviews the standard; it does not re-decide the case. A mechanism designed to reopen individual outcomes would fall outside the statistical-purposes gateway and would deserve to. This is the second occasion in this paper on which the law has already anticipated the seal (§3.2 is the first): the drafters were protecting individuals from being governed by research conducted about them, and in doing so they drew precisely the boundary between a review that can be lawful at scale and a review that cannot.

It follows that §5.2 and §5.3 rest on different bases and must not be run together. Cohort review is statistical processing about a standard. An individual's compelled review under §5.3 is an ordinary determination about that individual, on the institution's ordinary basis, with the ordinary rights attached. The two mechanisms meet only at the point where a pattern established by the first satisfies the threshold of the second.

The controller should not be the institution — but this must be constituted, not stipulated. The appointment condition (§5.4) has a consequence for data that is easy to miss: **the body the university may not appoint is also the body that should hold the linkage**. Where trajectory data must be joined to application records, the joining belongs in an independent environment, disclosing to the institution nothing but aggregate findings subject to ordinary disclosure control. But controllership under data protection law is not a matter of designation; it follows from who determines the purposes and means of processing, and it cannot simply be asserted as a safeguard. If the independent body determines why and how the linkage is analysed, it is the controller — or, where purposes are genuinely shared with the institution or a national linkage service, a joint controller with responsibilities allocated by arrangement. Achieving the intended separation therefore requires that the determining role actually be vested elsewhere, by the same statute or registration condition that supplies the appointing authority (§5.4), and not left to a data-processing agreement that a determined institution can shape. Constituted correctly, the arrangement does two things at once: it keeps the institution from acquiring a dossier on people who declined to remain in a relationship with it, and it removes the institution's capacity to identify — and so

to disadvantage — the individuals whose outcomes defeated its standard. Retaliation protection (§5.3) is a rule; this is an architecture, and where it can be constituted the architecture is the stronger of the two.

The counterfactual problem, and the asymmetry it forces. A selection standard operating under scarcity produces false negatives without being defective: where a hundred places are contested by a thousand able applicants, many who are rejected will succeed elsewhere, and their success shows only that the institution lacked infallible foresight, not that its criterion was unjustifiable when applied. Any evidential rule that treated later success among the excluded as *proof* of a defective standard would therefore prove almost every standard defective, which is to prove nothing. The mechanism must be more careful than that, and the earlier versions of this argument were not.

Two distinct difficulties have to be separated. The first is causal contamination: rejection is not a neutral observation of a candidate but an intervention upon them, so a rejected applicant who fares badly may have fared badly *because* they were rejected — evidence contaminated in exactly the direction that flatters the institution. The second is ordinary selection under scarcity, just described. The first makes poor outcomes among the excluded nearly worthless as evidence; the second makes good outcomes among the excluded insufficient, on their own, to condemn a standard. Together they force an asymmetry and a structure.

The asymmetry: **poor outcomes among the excluded are not evidence that the standard was sound.** A standard is not vindicated by the later struggles of those it turned away, because the turning away may have caused the struggle. This direction of inference is inadmissible, full stop.

The structure: good outcomes among the excluded are admissible, but they operate in three tiers rather than as a single defeater, and which tier is reached is a matter of degree and pre-specification, not of a single success story.

- **Tier one — trigger.** A credible instance of later success among the excluded is sufficient to *open* a review. It shifts nothing yet; it obliges the standard to be examined against its excluded population. This is a low bar by design, because triggering is not defeating.
- **Tier two — burden shift.** A *repeated, class-level* pattern of later success — not one striking case but a category of contribution the standard has declined across several cycles — shifts the burden of justification onto the standard. The institution must now show why the pattern is consistent with the standard's continuation, rather than the challenger showing why it is not.
- **Tier three — defeat.** The standard is defeated where the class-level pattern *materially exceeds a pre-specified expected false-negative rate*, or reveals systematic blindness to a type of contribution the standard cannot represent at all. Defeat requires the comparator: an account, fixed in advance, of how many false negatives a standard operating honestly under known scarcity would be expected to produce. A pattern within that expectation is the price of scarcity; a pattern beyond it is the signature of enclosure.

This is why **base rates are indispensable, not irrelevant** — the correction of an earlier and mistaken formulation. Without an expected false-negative rate, a comparator population, and an account of capacity constraints, there is no way to distinguish a defective standard from an imperfect standard doing its honest best under scarcity. The base rate is precisely what tier three measures against. What the mechanism denies is not the relevance of base rates but the sufficiency of averages: the question is never whether the excluded do worse *on average* — under scarcity they will — but whether a *particular class* of contribution, one the standard is structurally unable to see, exceeds the expected miss rate cycle after

cycle. Evaluative-layer capping predicts exactly that signature⁵⁶: the work that would reopen a standard is the work the standard misrecognises, so it surfaces among the excluded who succeeded in a *specific* class, not among the excluded at large.

A difficulty must be conceded here rather than concealed, because it presses on the comparator the whole tier depends on. The expected false-negative rate is calibrated against the population the standard *admitted* — those it judged correctly, or well enough — and that population is certified by the same enclosed standard whose blindness is in question. The denominator is, to that extent, contaminated in the way the numerator was: a standard that systematically misrecognises a class will also have mis-certified its own successes, so an expected miss rate derived from them understates the misses. This does not collapse the tier, for two reasons. The signature the tier looks for is a class-level *excess* that survives any plausible setting of the rate — if a category of contribution is declined cycle after cycle and vindicated elsewhere, no reasonable comparator absorbs it — so the mechanism can be built to require the excess to hold across a range of rates rather than a single point estimate. But the calibration cannot be made fully clean, and where the enclosure is deep enough to distort both admitted and excluded populations alike, the comparator degrades toward the recognition-order problem already stated above. The honest position is that tier three is robust against moderate contamination of its base rate and weakens as the contamination approaches totality — and that at the limit it is the plural, outcome-based evidence, not the rate, that carries the finding.

The recognition-order problem. One difficulty remains, and it is the deepest, because it threatens the evidence rather than the inference. ‘Success elsewhere’ is not a standard-free fact about a person. It is, almost always, recognition by *another* recognition system — a prestigious chair, a large grant, a citation record. A mechanism that defeats a standard whenever the excluded are later recognised elsewhere therefore does not escape the recognition order; it defers to a second node of it, and risks identifying only those whom *one* incumbent recognised after *another* failed to, while remaining as blind as the institution it audits to work that went unrecognised everywhere. The review would then reproduce, at one remove, the very enclosure it exists to detect (§1.2).

There is no clean escape, and the honest course is to state the limit and narrow the evidence rather than to pretend the problem away. The evidence admissible at tiers two and three must therefore be, so far as possible, **plural and not reducible to downstream institutional prestige**: demonstrated technical achievement, independently verifiable; uptake or use by communities the standard was supposed to serve; correction of an existing literature; durable practical application; generative effects — work that itself opened new work — that are not merely a later prize. Prestige-mediated recognition may corroborate, but it may not be the sole or principal defeater, precisely because it is the recognition order speaking in its own voice. This does not dissolve the problem; a review body drawing on ‘community uptake’ or ‘practical application’ is still exercising a standard, and cannot exhaustively validate it (§1.2). What it does is push the residual enclosure outward — from a single institution’s standard to the plural standards of several independent communities — which is improvement, not escape. The full treatment of a recognition order whose nodes are convertible into one another without any being answerable is the subject of separate work and is not resolved here.⁵⁷ What this paper claims is the narrower point: that later success is admissible evidence only as a *plural* signal, only in *class-level* patterns, and only against a *pre-specified* base rate — and that where even a plural signal comes back empty, the standard has not thereby been vindicated, because the enclosure may be shared.

⁵⁶ Kahl, ‘The Recognition Game’.

⁵⁷ Kahl, ‘The Recognition Game’; Kahl, ‘Conformance Is Not Constitution’.

5.3 Oversight Institutions: Standing to Trigger

Epistemic audit boards and fiduciary oversight committees are retained, on the four conditions set out in this section and the three that follow. The bodies the sector actually constitutes satisfy none of them. Without all four such a board is not merely weak but **worse than nothing**, because a body that cannot defeat a standard supplies the appearance that the standard has been tested.

Any student, member of faculty of any grade, member of staff, applicant, rejected researcher, or other party governed by a decision may **compel** its review. Not petition. Not request that the board consider whether to consider. Compel.

A route that no one is entitled to invoke is not available⁵⁸, and a board that reviews on its own initiative reviews what it has already decided to look at. This is the ordinary design: an audit body tasked with reviewing recognition practices, communications and curricular decisions periodically and on its own motion, with no beneficiary holding a key. Such a body examines the institution through the institution's own sense of what merits examination, which is the sense at issue.

The threshold is demanding, and it is **stated**. It is met where the applicant identifies (i) a decision and the standard said to have been applied to it, and (ii) either a specific outcome inconsistent with that standard as stated, or a class of contribution the standard has not recognised across a stated number of cycles.

The threshold cannot be, and is not claimed to be, *literally* non-discretionary. Whether an applicant has identified a genuine 'class of contribution' rather than an idiosyncratic grievance, whether a stated number of cycles has elapsed, whether an outcome is inconsistent with a standard as stated — each requires interpretation, and a scheme pretending otherwise would collapse at the first hard case. What the design excludes is not judgement but *the reviewed party's discretion over whether it is judged*. The distinction is between a threshold that is **bounded, reasoned and reviewable** — applied against stated criteria, with reasons recorded, and itself open to challenge before an independent tier — and a threshold applied *at large by the body it filters*. The first is ordinary adjudicative discretion, of the kind every tribunal exercises at the leave stage. The second is a veto wearing the clothes of a filter, and it will be exercised sincerely and still function as a veto, because the body cannot see the standard under which it exercises it. The requirement, precisely stated, is that the gatekeeping discretion not rest with the party whose standard the gate protects.

Note what the second limb does. It permits the sealed party to trigger review on a *pattern* rather than an argument. This is the concession the seal requires: the applicant is not asked to say what is wrong with the standard, because they cannot exhaustively, and neither can anyone else.

A public lawyer will object at this point that a stated, non-discretionary threshold courts the opposite vice. Administrative law treats the *fettering of discretion* as itself a ground of review: a decision-maker who applies a fixed rule mechanically, without genuine consideration of the case before it, acts unlawfully. Does a non-discretionary trigger not mandate the very fettering the law forbids?

It does not, and the distinction is the one that separates this proposal from an ordinary rule-bound scheme. The non-discretionary element is confined to the *trigger*: whether the stated conditions for opening a review are met is not left to the discretion of the body whose standard is under review, precisely because a discretionary gate held by that body is the delegated-discretion pathology in

⁵⁸ Kahl, 'Answerable Knowledge'.

miniature — authority to filter, exercised without evaluative control over what the filter excludes.⁵⁹ What happens *after* the trigger is not mechanical at all: the review itself is an exercise of judgement, reasoned and reviewable (§5.5), and the consequences that may follow (§5.12) are selected, on the evidence, by a body applying its mind to the case. Fettering is the abdication of judgement where judgement is owed; this proposal removes discretion only from the gatekeeping function, where discretion held by the reviewed party is not judgement but veto, and preserves it exactly where the law requires it — in the reasoned determination that follows. The threshold is bounded and rule-stated; the determination is bounded and reasoned. Neither is unfettered, and neither is fettered.

Retaliation against a party who triggers review — in supervision, assessment, promotion, renewal, reference, funding or reputation — is a disciplinary matter, and is stated to be one. Standing that costs the sealed party their position is not standing, and it is the condition most often left out of proposals that vest standing at all.

5.4 Oversight Institutions: Appointment

The board must not be appointed, convened, staffed, instructed, resourced or dissolved by the authority whose decisions it reviews.

This is not one safeguard among the fixed terms and transparent selection processes usually offered under the heading of independence. It is the condition on which all of them depend. Fixed terms in the gift of the audited party are fixed terms in the gift of the audited party. A transparent process by which the audited party selects its own auditor is a transparent capture.

Governance proposals in this area characteristically specify *how* appointments are made and never say *by whom*, which on the face of the text leaves the appointment with the audited institution's own governance. It is the largest hole in the standard design, and closing it is the greater part of what this chapter does.

What fails the condition can be stated exactly: appointment by council, senate, the vice-chancellor's office, or any committee those bodies constitute, staff or fund. What satisfies it is a question of charter and jurisdiction, and the answer will differ. Candidates:

- appointment by the sector regulator, where one exists and is not itself constituted by the regulated institutions;
- appointment by a standing cross-institution panel drawn from a register maintained outside any single institution, members allocated by lot and barred from reviewing their own;
- funding by a levy fixed *ex ante* and not adjustable by the institution within a review period;
- a dissolution power exercisable only by the appointing authority, on recorded grounds, and never during a live review.

If none is available — if the university appoints its own auditor — then the mechanism fails, and this paper says so rather than proceeding. **An unspecified appointing authority is not a gap to be filled at implementation. Unspecified means captured**, because the only body with the standing, the motive and the drafting resources to fill it is the audited party.

⁵⁹ Kahl, 'Delegated Discretion'.

Independence is not one thing, and the binary is too coarse. To speak of a body as ‘independent’ or ‘captured’, as the paragraphs above have, is a useful first cut and a misleading final one. A regulator may be structurally independent of any given university and still politically directed, funded at the discretion of a department of state, or cognitively dependent on the very metrics whose effects it would be reviewing. A cross-institution panel may be independent of the particular institution and wholly embedded in the same disciplinary and recognitional order — which, given the recognition-order limit at §5.2, is the independence that matters least and is hardest to obtain. External appointment secures one kind of independence and not the others. The condition should therefore be disaggregated, and a mechanism assessed on each axis rather than waved through on the first:

- **Appointment independence** — the reviewer is not chosen by the reviewed. This is the axis §5.4 has addressed, and it is necessary but not sufficient.
- **Budgetary independence** — the reviewer’s resources are not set, varied or withheld by the reviewed, whether directly or through a shared paymaster who shares the reviewed party’s interest. A body independently appointed but annually funded at the reviewed party’s discretion is captured on a lag.
- **Operational independence** — the reviewer controls its own agenda, timetable and staffing, and cannot be slowed, starved of information, or procedurally entangled by the reviewed party. Standing to trigger (§5.3) is worth nothing if the triggered review can be administered into abeyance.
- **Epistemic independence** — the reviewer does not apprehend the question through the same standard as the reviewed. This is the hardest, is never complete (§1.2), and is why the primary defeater is outcome-over-time rather than a better-qualified panel: an independently appointed, independently funded board that shares the enclosure evaluates the standard from inside it. Epistemic independence is approached not by finding unenclosed reviewers, who do not exist, but by admitting plural and outcome-based evidence that no single standard curates (§5.2).
- **Remedial independence** — the reviewer’s determinations take effect through a route the reviewed party does not control (§5.12), for a body that can find but not act is independent in its findings and captured in their consequences.

No appointing arrangement scores maximally on every axis, and the paper does not pretend one does. The claim is weaker and firmer: a mechanism weak on an axis should be known to be weak on that axis, so that reliance is placed elsewhere — most often on outcome-over-time, which is the one defeater that does not require the reviewer to be independent in a way no reviewer can be.

5.5 Oversight Institutions: Recorded Reasons for Refusal

Where a request to reopen is refused, the refusal and its grounds are recorded and published. The function of this must be stated, or the reader will take it for the transparency remedy §1.2 rules out.

It is not explanatory. A recorded reason cannot deliver the standard under which the refusal was made; it can only supply a rendering of that standard, produced under it. Its function is **evidential**. It fixes, at the time and before the outcome is known, what the board took the standard to require. It thereby builds the record against which later evidence can produce defeat under §5.2.

Refusals accumulate into a series. A series that subsequent events contradict is a defeated standard — and the contradiction is visible without anyone having had to represent the standard at all. This is the only work disclosure can do here, and it is enough.

5.6 Oversight Institutions: Composition

Composition should be diverse: junior and senior faculty, students, staff, community partners, external experts. The usual claim made for that diversity should not be.

Diverse composition improves **epistemic weight**. It widens the range of considerations a body can bring to bear, and that is a real good, worth having for its own sake. It does not protect against capture, and it does not prevent finality. A nine-member board can place a conclusion beyond reopening as completely as one person can — and when it does so, it does so with the *appearance* of having already been scrutinised, which is worse than the appearance of not having been, because it forecloses the suspicion that would otherwise attach.

*A plural body that finalises is the pathology; a plural body that reopens is the remedy; the plurality is incidental to both.*⁶⁰ What prevents finality is not the number of hands but the fact that the capacity to reopen is not controlled by the authority whose conclusion would be reopened.

The error here is precise, and worth naming because the whole sector makes it: epistemic weight is conflated with epistemic finality, and the distribution of the former is treated as though it dissolved the latter. It does not. It relocates finality to a committee and gives it a quorum.

5.7 The Regress Objection

Specifying an appointing authority invites the oldest question in the literature. *Quis custodiet ipsos custodes?* Someone appoints the appointer, and someone appoints them.

The answer is neither a better appointer nor a more plural one. Both replies assume the regress is a series that must terminate somewhere, and that the task is to locate the node at which it stops. There is no series.

The seal attaches to the notion **locus of closure as such**.⁶¹ Every locus of closure is a generative condition for the claims that pass through it. Every one is therefore sealed both from what it governs and from itself. Therefore every one must be defeatable. This is one theorem applied at every node — and the audit board is a node like any other, as are the regulator, the cross-institution panel, and this paper.

So the board is not required to be appeal-proof, and no appointing authority is required to stand beyond question. The board is required to be **non-self-controlled** — it does not hold what it reviews — and to be transparent as to its own determinations in the evidential sense of §5.5: thresholds stated in advance, reasons recorded, review dates fixed. Its determinations are then open to a successor panel, to the wider institution, and above all to later evidence under §5.2.

The regress does not need to stop. It needs to stay open at every point. That is a different demand, and unlike termination it is satisfiable.

⁶⁰ Kahl, 'Answerable Knowledge'.

⁶¹ Kahl, 'The Inside of Generatedness' §6.2.

5.8 Decentralised Credit and Recognition Infrastructure

Traditional systems of academic reward privilege visible outputs — rankings, publications, high-profile grants — at the expense of less visible but vital epistemic labour. An alternative credit infrastructure should redistribute recognition to encompass mentoring, collaborative teaching, community scholarship and technical contribution. This aligns with the pedagogy of collective empowerment⁶² and operationalises Anderson’s epistemic democracy.⁶³

Under §5.1 its status must be stated honestly. This is a mechanism of **epistemic weight**, not of answerability. It changes who is recognised; it does not make any recognition standard capable of losing. It therefore does not carry the remedy and must not be presented as though it did.

It earns its place for a different reason. A credit record that registers what was contributed, by whom, and when — recorded contemporaneously, not reconstructed at review — is a principal source of the outcome data on which §5.2 depends. Like recorded refusals, its value is evidential, and for the same reason.

It is subject to the same four tests as the boards. *Standing*: any contributor, named or omitted, may compel correction of the record. *Appointment*: the ledger is not maintained by the body whose recognition decisions the ledger evidences. *Plurality*: a plural committee curating credit is a plural committee curating credit. *Legitimacy*: a credit infrastructure adopted because it displays well is hollowed at founding.

5.9 Visual Plurality Protocols

To counteract epistemic inversion, institutions should adopt protocols governing official communications, so that announcements of achievement show the range of contributors rather than senior leadership alone.⁶⁴

These protocols are cosmetic in the strict sense: they govern the optics. That is not an objection — epistemic inversion is a pathology of optics and is properly met at the level of optics. But they make no standard capable of losing, and they carry a specific exposure the rest of this chapter does not: **the visual display of plurality is precisely the artefact an optocratic institution is best equipped to produce**. A photograph of a diverse team is producible without any change whatever to who is recognised, promoted, funded or heard.

The protocols are therefore retained as hygiene, and nothing in them discharges fiduciary answerability. And I record the prediction of §3.6 again in its sharpest form, because this is where it will be tested: of all the mechanisms proposed here, visual plurality protocols are the one most likely to be adopted enthusiastically and alone. Where that happens, the adoption is not partial progress. It is the diagnosis, confirmed.

⁶² Freire, *Pedagogy of the Oppressed*; Darder, *Reinventing Paulo Freire*.

⁶³ Anderson, ‘The Epistemology of Democracy’.

⁶⁴ Barnett, *Realizing the University*.

5.10 Stakeholder Engagement

Open forums, deliberative assemblies and consultation platforms are mechanisms of **voice**. Voice is not standing.

The standard call is for continuous consultation, which describes a channel the institution opens, defines, schedules, staffs and may close. Nothing within it can compel anything. Consultation that the consulted cannot convert into a review the institution must conduct is the form of scrutiny retained around conclusions the form may not reach: **hollowed authorisation**⁶⁵, and the sector's most practised product.

The forums are retained for what they actually do, which is not nothing. They are how a standard's costs become sayable, and how a party who cannot represent the standard can nonetheless report what living under it is like. That reporting is evidence for §5.2 and material for the pattern limb of §5.3.

But they are re-routed. An issue raised in a forum that meets the stated threshold triggers review **as of right** — without the forum's convenor, the institution, or the board deciding whether it does. A forum that feeds a mandatory pathway is a fiduciary instrument. A forum that feeds a report is a suggestion box with better furniture.

5.11 Phased Implementation and the Order of Phasing

Implementation may be gradual. What may not be phased is the check.

Pilots, phased roll-outs and iterative revision are ordinarily proposed with insulation from capture listed alongside them as a safeguard. The order is wrong. The appointment condition and standing to trigger come *first*, because a pilot that lacks them is not an early version of this framework. It is a pilot of the thing this framework replaces, and it will report success.

A sequence consistent with the framework:

1. Fix the appointing authority and the funding basis, in charter or by regulation, before any board exists.
2. Vest standing, with the threshold stated and retaliation protection in force.
3. Begin contemporaneous recording — refusals, credit, and above all the outcomes of the excluded — from the first day, since §5.2 cannot review a record that was never kept.
4. Set the first review date at adoption, and do not move it.
5. Extend across faculties.

There is a reason the appointment condition must be fixed at sector level, and §5.4 understates it by presenting independence solely as a guarantee of non-self-control. The degraded order is an equilibrium⁶⁶: it persists not because it is preferred but because the first institution to leave it bears the cost alone — publishing defeats while its competitors publish assurances, into a market that reads the difference as weakness rather than as candour. Independent appointment is therefore also the

⁶⁵ Kahl, 'Answerable Knowledge'.

⁶⁶ Kahl, 'The Recognition Game'.

coordinating device that makes the better equilibrium safe to reach. An institution acting unilaterally has the courage and not the device, and courage is not a mechanism (§7.3).

One criterion of pilot evaluation is ruled out. A pilot assessed on stakeholder satisfaction, staff morale or reputational effect is assessing whether the mechanism made people comfortable, which is very nearly the opposite of whether it worked. The criterion is whether any standard was defeated — and if none was, whether the record shows why not.

5.12 What Follows When a Standard Loses

The mechanisms to this point establish who may trigger review, who must not control it, and what evidence defeats a standard. They are silent on the question that decides whether any of it matters: what *follows* when a standard is defeated.

The silence is not incidental, and it is the point at which this chapter is most exposed to its own objection. On the construction of breach given at §1.1, an evaluative standard maintained without answerable defeasibility is *already* in breach, independently of any bad outcome; the consequences specified here are the remedy for a breach that consists in the standard's condition, not a punishment for a discovered wrong. A finding of defeat that carries no consequence is a published finding, and a published finding is exactly the thing the institution has shown it can absorb — an adverse report metabolised into evidence that the institution takes scrutiny seriously. A review body that can find against a standard but can do nothing about it is a review body that supplies assurance by another route: the assurance that findings are made. **Defeat without a consequence is the audit trap (§5.14) in the remedy's disguise.**

So the framework must specify a law of consequences, and it must do so in advance, because a consequence settled after the finding is a consequence the losing party helps to choose. What follows is not a single remedy but a defined set, selected against the finding by a rule fixed at adoption.

On a finding of defeat, the review body may:

1. **Suspend** the standard, so that it may not be applied pending reconstruction. This is the default where the finding is that the standard is producing the exclusion, rather than that a better standard is known.
2. **Remit** the standard to its constitutor for reconstruction, with reasons, within a fixed period — the constitutor being whoever in fact holds the power to reconstitute it, which under §5.13 is frequently not the reviewed institution.
3. **Substitute** a specified default standard in the interim, where one exists and is less exclusionary on the evidence.
4. **Reopen affected decisions**, prospectively or, where feasible and proportionate, for the identifiable class the standard excluded — subject to the sharp limit below.
5. **Declare** the defeat without individual remedy, where the standard's harm is systemic and no individual reopening is proportionate, the declaration itself carrying the consequences at (1)–(3).

Two limits are structural rather than prudential.

Reopening is bounded by the seal, not merely by cost. A review that defeats an admissions standard does not thereby establish that any particular rejected applicant should have been admitted — the standard's defeat is a fact about the standard, not about the counterfactual merits of a person, which no one is positioned to reconstruct. Individual reopening is therefore available for *procedural* restoration (the applicant is admitted to the next competition under a reconstructed standard) and not for *substantive* reversal (the applicant is treated as though they had been admitted). This mirrors the academic-judgment limit at §1.1 rather than fighting it: the mechanism reopens the standard, not the score.

The consequence must not be defeasible by the loser. An institution that may decline to suspend, or may set the reconstruction period, or may choose whether to remit, holds the capacity to reopen its own conclusion — which §1.2 forbids. The selection among (1)–(5) belongs to the review body under the rule fixed at adoption, not to the institution whose standard was defeated. Where the institution controls the consequence, the finding is advisory, and an advisory finding is the assurance-producing artefact this paper exists to refuse.

There is a cost to stating consequences this plainly, and it should be named. A framework that can suspend a standard, remit it, and reopen a competition is a framework with teeth, and teeth are what institutions defund. The resourcing and appointment conditions of §5.4 are therefore not administrative preconditions of the consequences but structural ones: a consequence the institution can starve is a consequence the institution controls, one step removed.

5.13 Where the Standard Was Not Constituted by the Institution Applying It

The framework so far assumes that the institution holding a standard can reconstruct it. Frequently it cannot, and the assumption is the last place the seal is allowed to hide.

Many of the standards by which a university recognises and excludes are not constituted by the university at all. Research-assessment frameworks, professional-body accreditation, funder eligibility rules, league-table methodologies, immigration requirements attached to sponsorship, and sector-wide metric conventions are constituted elsewhere and *transmitted* by the university, which applies them under external penalty for departure. A department applying a REF-shaped notion of quality, or a programme applying an accreditor's curricular constraint, is an operator of a standard it did not write and cannot unilaterally abandon.

For this class the distinction between three things that the framework has so far run together becomes decisive: the **constitution** of a standard, the **answerability** for it, and the **standing** to demand an account are independent, and a provenance trace connecting them requires a stopping rule.⁶⁷ It follows that a review body which terminates its inquiry at the university applying the standard will hold a local operator answerable for a rule constituted at sector level — which is neither just to the operator nor effective against the standard, since the operator has no power to reconstitute what it received. This is delegated discretion in its administrative-law form: authority to apply a standard has been distributed down the chain to a party that exercises it without evaluative control over it, while the control sits — if it sits anywhere — with a constitutor the applied-to party cannot reach.⁶⁸ The structure is familiar from

⁶⁷ Peter Kahl, 'Conformance Is Not Constitution: Provenance, Answerability, and Standing' (Lex et Ratio Ltd 2026) <<https://doi.org/10.13140/RG.2.2.28710.28486>>.

⁶⁸ Kahl, 'Delegated Discretion'.

the law of agency, where an agent exercises delegated authority under instructions whose interpretation is itself constrained by the principal's purposes, and where the agent's fiduciary character does not dissolve merely because the evaluative direction originates elsewhere.⁶⁹ A university transmitting an externally constituted standard is, in this respect, an agent applying instructions it did not write; the answerability for the instruction's content does not terminate at the agent. The attribution provision is the remedial counterpart: it directs the finding to wherever the evaluative control actually rests.

The framework therefore requires a **standard-attribution provision**:

Where the review finds that the operative standard was constituted, wholly or in part, outside the institution applying it, the finding must travel to the constitutor with jurisdiction and remedial competence over the standard, rather than terminate at the point of application.

Three consequences follow.

First, the consequences of §5.12 attach *at the locus of constitution*. Suspension, reconstruction and substitution are demanded of whoever can actually perform them. Against a transmitting university the available consequence is narrower and specific: a duty to record and to transmit the finding upward, and — where the external standard is defeated but not yet reconstituted — a defence to the charge that the university failed in its own answerability by continuing to apply what it had no power to change, provided it has transmitted and recorded.

Second, this closes an evasion the single-institution framework leaves open. Without attribution, every institution can truthfully say that the standard is not really its own — that it is only doing what the REF, the accreditor or the funder requires — and the finding dissipates across a chain in which no link is answerable. Attribution converts that from an excuse into a routing instruction: not *the standard is not ours, so there is nothing to answer*, but *the standard is not ours, so the finding travels to whoever's it is*.

Third, it exposes a limit of this paper, which is stated rather than concealed. Where a standard is co-constituted across bodies that are themselves recognition orders — a funder deferring to metrics that defer to journals that defer to the funder — attribution identifies a circle rather than a terminus. The stopping rule for such traces, and the constitutional defect of a recognition system whose standards are convertible into one another without any of them being answerable, is the subject of separate work and is not resolved here.⁷⁰ What this paper claims is narrower: that the finding must travel, that it must travel to a locus with remedial competence, and that where it reaches only a circle, the circle is itself the object requiring constitution — which is a finding, not a failure.

5.14 Legitimacy and the Audit Trap

Trust in institutions is in global decline, and the temptation to offer the mechanisms of this chapter as the answer to that decline must be refused. It is not a small temptation, and it is the one to which accountability proposals almost invariably yield: audit boards, accountability reports and stakeholder forums are offered as tools of legitimacy, and commended as a model for restoring public confidence. That framing is fatal, and nothing here should be read as making it.

⁶⁹ Deborah A DeMott, 'The Fiduciary Character of Agency and the Interpretation of Instructions' in Andrew S Gold and Paul B Miller (eds), *Philosophical Foundations of Fiduciary Law* (OUP 2014) 321.

⁷⁰ Kahl, 'The Recognition Game'; Kahl, 'Conformance Is Not Constitution'.

A body designed to produce legitimacy is designed to ratify. Its output is trust; trust is what it is for; and an instrument that exists to produce trust cannot withhold it, because withholding it is failure by the instrument's own criterion. Its adverse findings are its malfunctions. Over a few cycles, an institution that has commissioned an assurance-producing body and received assurance will conclude that it has been audited, and the conclusion will be sincere, and there will be a report.

This is Power's *Audit Society* thesis exactly⁷¹: audit expands not because it verifies but because it produces comfort about verification, and its rituals substitute for the thing they are held out to check. It is the thesis Chapter 3 deploys against everyone else's compliance apparatus — against diversity statements and accreditation claims produced for external optics — and which any accountability proposal walks into the moment it promises confidence as its product. **Audit that produces trust is audit that has stopped being able to withhold it.**

A companion mechanism should be named here, because a board under pressure will reach for it within a cycle. Where a determination is difficult to defend, the metric is promoted from check to verdict, and the answerability for the misrecognition is **laundered**: no one decided, the number decided.⁷² This is why §5.5 requires recorded *reasons* and not a recorded *result*. A board that answers a compelled review by reciting a score has not reviewed anything. It has relocated finality into an instrument that cannot be argued with, and it has done so while producing a document that looks exactly like accountability.

The correct statement is the inverse, and it is uncomfortable enough that it should be made plainly.

A board whose function is to generate trust cannot perform the function proposed for it in Chapter 5. And a board that performs the function proposed for it in Chapter 5 **will frequently reduce the institution's legitimacy — which is the point.** It will find that recognition standards excluded contributors who should not have been excluded. It will record that. The record will be public, and it will be damaging, and the damage is the mechanism working. An institution that establishes such a board for reputational reasons has established a board it will defund at its first real finding — which is why the resourcing condition at §5.4 is not administrative detail but structural.

Whether public trust follows is not the mechanism's concern, and cannot be made its concern without destroying it. Trust may follow, in time; there is a decent argument that an institution known to hold defeasible standards would eventually be trusted more, and more justifiably, than one that publishes assurances. That argument is **out of scope**, and any attempt to bring it into scope is fiducial hollowing performed on the remedy itself.

The reason is fiduciary rather than tactical, and it is the oldest reason in this paper. The duty runs to the sealed party: the student, the applicant, the junior colleague, the rejected researcher who bears the standard's costs and can neither monitor nor replicate the deliberation that produced it.⁷³ It does not run to the public's opinion of the institution. Where the two diverge — and they will diverge exactly when the board finds something — the beneficiary wins.

That is not a hard case. It is what a fiduciary duty *is*. *Keech* settled the point three centuries ago against a trustee who had taken a benefit no one could show had harmed anybody, and the rule was applied not

⁷¹ Michael Power, *The Audit Society: Rituals of Verification* (OUP 1997).

⁷² Kahl, 'The Recognition Game'.

⁷³ Kahl, 'Epistemic Humility as Fiduciary Obligation'.

in the least relaxed. An institution that reserves the right to weigh its beneficiaries' interests against its own reputation has not misapplied fiduciary law. **It has declined it.**

6. Anticipating Resistance and Counter-Arguments

Every governance reform generates resistance. Epistemocracy will be no exception: it reorders credibility, redistributes recognition, and makes standards defeasible. This chapter anticipates the likely lines of opposition. Four come from outside the framework. The fifth, at §6.5, is the only serious one.

6.1 Resistance from Senior Administrators

Under optocratic regimes, credibility is centralised in the figure of the Vice-Chancellor, President or Rector. Press releases, brochures and ceremonial speeches allow administrators to stand as proxies for institutional excellence.

Objection. Distributing credibility across the institution, and subjecting recognition standards to compelled review, will diminish the authority of senior leaders, weaken external fundraising, and confuse accountability lines.

Response. Epistemocracy does not abolish leadership; it re-grounds it. Note first what it does *not* do: it does not diffuse procedural authority (§1.2). Someone must convene, decide the order of business, and be identifiable as having done so. Diffusing that produces not accountability but the absence of anyone to hold, and the framework concentrates it deliberately. What the framework removes is not authority but *finality* — the standing to place a conclusion beyond reopening because of who reached it. No account of leadership that survives scrutiny requires that standing.

The fundraising limb should be answered directly rather than accommodated. It asks the institution to weigh its beneficiaries' interests against its revenue from third parties, and a fiduciary does not have that permission. Whether legitimacy or donations follow is addressed at §5.12, and the answer there is that they frequently will not.

Strategic framing. Administrators who claim their roles will be devalued by the loss of finality are confessing that their authority rested on finality. That is not a defence. It is the diagnosis, offered in the first person.

6.2 Resistance from Managerial Culture

Managerialism thrives on measurability. Rankings, KPIs and branding exercises create the appearance of accountability while hollowing out fiduciary duty.⁷⁴

Objection. Opacity-respecting pluralism is impractical. How can a university govern without standard metrics? Qualitative knowledge cannot be managed or compared. Regulators and funders demand measurable outputs.

Response. The objection rests on a false binary: measurable versus ungovernable. Institutions can adopt plural metrics — narrative reporting, qualitative evaluation, peer recognition structures. Biesta reminds

⁷⁴ Slaughter and Rhoades, *Academic Capitalism*.

us that genuine education entails risk and unpredictability⁷⁵; to erase uncertainty by reducing all knowledge to numbers is not governance but abdication of the educational task.

But the deeper reply is that §5.2 is *more* empirical than what it replaces, not less. Periodic review against later evidence is measurement — of outcomes, over time, including the outcomes of the excluded. The managerial objection is not that epistemocracy is unmeasurable. It is that epistemocracy measures the wrong population. Current metrics measure the institution's successes; §5.2 measures its rejections. There is nothing soft about that, and nothing about it that a regulator could not require.

One clarification forecloses a misreading that would otherwise be fatal. **Epistemocracy is not anti-metric, and cannot be** — because metrics were themselves a reform, and a real one, against unaccountable discretion. The corrective to metric governance is therefore not a return to discretion; that is the arrangement metrics were introduced to end, and it is remembered fondly by exactly the parties whose discretion it protected. What epistemocracy asks is that the metric operate as a **check** rather than a **verdict**⁷⁶: available to defeat a judgement, never able to constitute one, and never able to absorb the answerability for it (§5.12). Judgement is owned, named, and answerable; the number is evidence against it, not a substitute for it.

Strategic framing. Managerial objections reveal a dependency on ranking systems that operate as external instruments of control while being described as internal instruments of quality. Exposing the dependency is itself part of the case.

6.3 Resistance Based on Resource Scarcity

Objection. We lack the resources. Budgets are strained by government cuts, student demand and global competition. Audit boards, credit infrastructures and independent appointment are luxuries.

Response. The objection presumes the duty is discretionary. It is not: a fiduciary does not plead expense against loyalty. Where resources genuinely cannot support the mechanisms of Chapter 5, the honest fiduciary statement is that the institution cannot presently discharge the duty — not that the duty has been scaled to the budget. The second is what will be said, and it is what fiducial hollowing sounds like in a finance committee.

Some phasing is nonetheless possible, in the order set at §5.11 and no other. Contemporaneous recording is nearly free and must start immediately, since §5.2 cannot review a record that was never kept. Nor is the evidentiary base the expense it is assumed to be: the sector already funds and operates outcome-tracking machinery, and §5.2 asks that it be pointed at a different population rather than rebuilt. Distributed credibility requires no new allocation; it is a change to promotion, tenure and recognition criteria. The genuinely costly item is independence — and that cost is not incidental, because independence is what the institution is being asked to buy.

Strategic framing. The appeal to scarcity is itself a symptom of academic capitalism⁷⁷: resources are found for optics and not for exposure. Reallocation, not expansion, is the fiduciary response — and the first place to look is the marketing budget, which is the drift, capitalised.

⁷⁵ Biesta, *The Beautiful Risk of Education*.

⁷⁶ Kahl, 'The Recognition Game'.

⁷⁷ Slaughter and Rhoades, *Academic Capitalism*.

6.4 Resistance from Inertia and Cultural Norms

Objection. Epistocracy is too disruptive. Universities are already under pressure; destabilising existing norms will create confusion and undermine morale.

Response. This objection illustrates the pathology it defends. Institutions mistake inertia for stability. Clinging to hierarchical optics undermines resilience and leaves universities brittle in the face of external shocks; standards that have never been tested are not robust, merely untested. Barnett's supercomplexity is the relevant frame⁷⁸: an institution whose function is to help a society live with irreducible conflict cannot itself require internal consensus to operate.

The morale limb deserves an honest answer rather than a reassuring one. Morale will fall, at first, among those whose standards are defeated — and §5.12 says why that is not an argument. It should not fall among those the standards excluded, who have been living with the outcome all along without the vocabulary to name it.

Strategic framing. Cultural inertia is not neutral; it is vested interest under the description of tradition.

6.5 The Objection from Within the Framework

The serious objection is the paper's own.

This paper is a locus of closure. It advances an evaluative standard — the test at §5.1 — under which certain mechanisms count as remedies and others as decoration. By its own theorem, I cannot exhaustively validate that standard from within: it is the law under which my own representing here proceeds. The paper therefore cannot certify itself, and its author is the last party who should be asked whether it is right.

I accept this without reservation. It is not a difficulty for the argument; it is the argument, applied where it applies. What follows from it is not that the paper is unfounded, but that the paper owes what it demands of everyone else: **it must be capable of losing**, and I must say in advance what would defeat it, before I know the outcome, so that the saying is evidence rather than rationalisation (§5.5).

The paper is defeated if:

1. **The seal is answered in its narrowed form.** The paper does not stake itself on the strong claim that an institution can represent none of its criteria; it grants that reflection reaches some. It stakes itself on the narrow claim that reflection cannot be *exhaustive and self-validating* — that no authority can occupy a standpoint outside its own standard from which the standard's full contribution is visible and can be certified correct. If it can be shown that some authority does obtain such a standpoint, so that diligent reflection and disclosure are in principle sufficient to validate a standard from within, then transparency is a sufficient remedy after all, §4.5 collapses, and with it the case for making exposure primary rather than supplementary. (Showing merely that institutions can identify *particular* criteria does not answer the seal; the narrowed claim already concedes that.)

⁷⁸ Barnett, *Realizing the University*.

2. **Plurality does the work I deny it.** If plural bodies are found reliably to reopen conclusions against the interests of the authority that appointed them, at rates comparable to non-self-controlled bodies, then §5.6 is wrong and the demotion of plurality was a mistake.
3. **The comparator cannot be cleaned enough.** Tier three (§5.2) requires an expected false-negative rate calibrated against a standard's admitted population, which the standard itself certified. If it can be shown that this contamination is not merely moderate but pervasive across realistic cases — so that no setting of the rate, and no plural non-prestige evidence, can distinguish a defective standard from an honest one under scarcity — then the primary mechanism cannot in practice discriminate the cases it exists to discriminate, and the paper's central evidential claim fails.
4. **The mechanisms run empty.** If an institution implements Chapter 5 fully — appointment condition, standing, contemporaneous recording, fixed review — and over several cycles no standard is ever defeated, *while* cohort data on the excluded show the exclusion patterns Chapter 2 predicts, then the mechanism does not do what I claim, and its adoption has cost the sector money and supplied it with assurance, which is the specific harm §5.12 warns against.
5. **The prediction fails, in the wrong direction.** §3.6 and §5.9 predict that vocabulary and visual protocols will be adopted and mechanisms will not. If institutions adopt the mechanisms readily, my account of institutional incentive is wrong and the diagnosis is weaker than Chapter 3 claims. I should be delighted to lose on this one.

Each of these is an outcome, not an argument, and that is deliberate: they are the kind of thing that can defeat a standard whose holder cannot see it. I do not have the view of this paper. Nobody does. What I can do is fix the dates and record the reasons, and that is what §§3.6, 5.9 and this section do.

Conclusion to Chapter 6

Resistance will be fierce because the proposal targets entrenched privileges: the optics of senior administrators, the security of managerial metrics, the comfort of scarcity arguments, the inertia of tradition. Each of the first four lines of resistance is itself evidence of pathology, and the fiduciary frame supplies a consistent reply: universities are not free to privilege optics, austerity or inertia over duties owed to those who bear the costs of their standards.

The fifth line is not resistance but the framework working, and it is answered in the only currency the framework accepts: stated defeat conditions, recorded in advance.

7. Theoretical Depth and Further Implications

7.1 Agonistic Pluralism: What Conflict Can and Cannot Do

Mouffe insists that democracy thrives not by suppressing conflict but by rendering it productive.⁷⁹ Agonistic pluralism conceptualises dissent not as an obstacle to democratic governance but as constitutive of it: a thriving democratic culture embraces disagreement and contestation rather than eliminating them.

⁷⁹ Mouffe, *Agonistics*.

It is often held that audit boards, visual protocols and stakeholder deliberations institutionalise structured dissent, and that the structured dissent is itself the safeguard. It is not. Contestation an institution values, hosts, schedules and finds productive is contestation the institution controls, and a standard is not endangered by disagreement its holder has convened. Structured dissent is a plural body's activity; it bears on finality only if the dissent can *compel* something. Hosted contestation characteristically compels nothing.

What agonism supplies instead is worth having on its own terms: an account of what a live reopening pathway feels like from the inside. An institution whose standards are genuinely defeasible will be uncomfortable, contested and short of consensus, and it should not read those symptoms as failure — nor should its regulators, its staff surveys, or its press office. Agonism is the phenomenology of the remedy. It is not the remedy, and the difference is that one of them can be produced by holding a forum.

7.2 Risk and Uncertainty as Education's Essence

Biesta argues that education is essentially an encounter with risk and unpredictability rather than the safe delivery of predetermined outcomes.⁸⁰ Under optocratic regimes, risk is domesticated — repackaged as marketing or controlled optics.

Epistemocracy reorients institutions toward educational risk as a fiduciary–epistemic duty, and adds a structural reason for it that Biesta does not need but the present argument does. Risk is not merely educationally valuable; it is *epistemically load-bearing*. An institution that eliminates unpredictable outcomes eliminates the only evidence that can defeat its standards. Under the seal, the surprise is not a cost of the system. It is the signal (§5.2), and the institution that has stopped being surprised has not stopped being wrong.

7.3 Virtue Epistemology and Its Limits

Zagzebski's account of intellectual virtue — courage, integrity, justice — supplements fiduciary theory.⁸¹ Fiduciary obligations are not merely procedural rules; their good exercise requires cultivated dispositions. In 'Directors' Epistemic Duties'⁸² I argued that epistemic openness demands more than structural compliance: it requires courage in disclosing uncomfortable truths, integrity in resisting performative gesture, and justice in recognising marginalised contributors.

All of that stands, within its scope, and the scope is narrower than a virtue-based account of governance requires. Distributed credibility embodies justice; opacity-respecting pluralism requires courage; both are exercises of epistemic weight. But virtue cannot carry answerability, and this is where the traditions of §4.2 divide most sharply.

A virtue is a disposition of a knower toward what the knower can see. The seal concerns the residue no knower reaches: not their whole standard, some of which reflection does recover, but the part that recovery is conducted under and cannot itself reach. A courageous holder of a hollowed standard discloses, courageously, everything they can see — and the residue is precisely what they cannot. A humble one is humble about the wrong things — which is what humility, unaccompanied by a

⁸⁰ Biesta, *The Beautiful Risk of Education*.

⁸¹ Zagzebski, *Virtues of the Mind*.

⁸² Kahl, 'Directors' Epistemic Duties'.

mechanism, reliably is. Fricker's virtuous hearer corrects for prejudice they have identified⁸³; the exclusions of §2.2 are the ones they have not.

The framework therefore asks nothing of anyone's character. **This is a feature.** A remedy that requires virtuous office-holders has selected for the institutions that least need it, and is unavailable exactly where it is wanted.

There is a second route to the same conclusion, and it does not pass through the seal at all. Grant, for argument, a knower who *could* see their own standard. The recognition game shows they still could not act on it alone: the arrangement is an assurance equilibrium, and unilateral departure is punished by the parties who control advancement.⁸⁴ Virtue fails here twice over — once because it cannot see, and once because seeing would not be enough. The arguments are independent, and a critic who wants virtue back must defeat both.

7.4 The Ecological University and the Question of Who Can Compel

Barnett's later work supplies the most developed alternative vision of the university's outward duties, and it is worth setting beside epistemocracy because the two agree substantially about the diagnosis and differ instructively about the remedy.

Barnett's ecological university is answerable outward: to knowledge, to learning, to persons, to culture, to social institutions, to the economy, and to the natural environment — a set of ecosystems the university depends upon and does not command, and whose demands upon it do not reconcile. His most recent statement makes the irreconcilability central, setting out eight ecosystems together with their *antagonisms*, and issuing a manifesto rather than a model.⁸⁵ This is the mature form of the supercomplexity thesis⁸⁶ and of the philosophical programme he has done more than anyone to constitute as a field.⁸⁷

Two things follow for the present argument.

The first is agreement, and it is deep. A university answerable to antagonistic ecosystems cannot optimise; it cannot possess a standard that reconciles its obligations; and it therefore cannot be in the position of knowing that its evaluations are correct. Barnett arrives from the direction of philosophical anthropology at something structurally close to §1.2: the university does not have the view.

The second is where epistemocracy departs, and the departure is narrow but consequential. Barnett's ecosystems have claims on the university; they do not have *standing* against it. Nothing in the ecological framework identifies who may compel the university to reopen a determination, on what threshold, at whose appointment, or with what consequence for refusal. The manifesto addresses the university's self-understanding — it asks the institution to conceive itself differently, and to feel the antagonisms. Under §5.1 that is a mechanism for making someone careful.

⁸³ Fricker, *Epistemic Injustice*.

⁸⁴ Kahl, 'The Recognition Game'.

⁸⁵ Ronald Barnett, *Realizing the Ecological University: Eight Ecosystems, Their Antagonisms and a Manifesto* (Bloomsbury 2025).

⁸⁶ Barnett, *Realizing the University*.

⁸⁷ Ronald Barnett, *The Philosophy of Higher Education: A Critical Introduction* (Routledge 2022).

I do not offer this as a refutation, because it is not one: Barnett is answering a different question, and answering it far more richly than epistemocracy attempts. But the questions meet. If the university genuinely cannot reconcile the demands upon it, then its determinations are permanently provisional, and a permanently provisional determination held by an authority that controls whether it is reopened is simply finality with better manners. **The ecological university needs a fuse.** What epistemocracy contributes to Barnett's programme is not a better vision but a smaller and more mechanical question: *who, precisely, may compel this university to look again, and what happens when it declines?*

7.5 Sustainable Epistemic Cultures

Barnett describes the supercomplex university as one that does not resolve complexity but learns to live with it.⁸⁸ Epistemocracy cultivates that capacity — not by embedding pluralism and transparency, which is the usual prescription, but by removing the institution's power to close a question permanently. A culture in which standards are known to be defeasible is a culture in which the unsettled state of a question is not an embarrassment requiring management.

Such cultures are resilient against populism, market instrumentalism and authoritarian knowledge control for a reason worth stating precisely: each of those is, in the terms of §1.2, a bid for **epistemic finality** — the standing to place a conclusion beyond reopening because of who holds it. The populist claims it for the people, the market for the price, the authoritarian for the office. An institution whose own architecture denies that standing to itself has the only principled ground on which to deny it to them.

8. Beyond the University: Transferability and Its Conditions

The pathologies identified in academia are not confined to universities. Optocratic drift, fiducial hollowing and epistemic inversion recur across corporate governance, NGOs and democratic polities, and so does the remedy — on conditions that §8.4 states, and that accounts of transferable governance models routinely omit.

8.1 Three Domains, One Structure

The pathologies recur beyond the university with their sealed parties differently placed, and in each the missing element is the same. In **corporations**, the fiduciary vocabulary is already present — directors owe care and loyalty under the Companies Act 2006 and at common law — and only the epistemic limb is absent; what epistemocracy adds is audit answerable to someone other than the appointing board, standing for employees and whistleblowers to *compel* review rather than report to a hotline the company runs, and periodic review of the board's risk standards against what later happened, with the excluded cases in the sample. Smith's analysis applies without adjustment⁸⁹: a power held for other-regarding ends may not be exercised in a conflict of duty and interest, and a board controlling whether its own determinations are reopened stands in precisely that conflict — permanently, structurally, and so universally that the universality has been mistaken for acceptability. In **NGOs**, the sealed party is unusually stark: a beneficiary population that bears the cost of programme standards, cannot monitor

⁸⁸ Barnett, *Realizing the University*.

⁸⁹ Smith, 'The Law of Loyalty'.

their deliberation, and, unlike a shareholder or student, typically has no exit; the structural blind spot is that impact reporting counts those served while the excluded appear in no denominator. In **democratic institutions**, political optics collapse plural deliberation into images of leaders — epistemic inversion (§3.3) on a different constituency — and what such institutions require is what universities require: a route to compelled reopening the decision-maker does not control, which judicial review, free elections and a free press already imperfectly attempt. The fiduciary–epistemic frame states in one sentence what those institutions are *for*: they exist because no authority can validate its own evaluative standard, and so none may control whether its conclusions are reopened.

8.4 What Transfers, and What Does Not

Two things do not transfer, though accounts of transferable governance models routinely imply that they do.

The vocabulary transfers too easily. Optocratic drift, fiducial hollowing and epistemic inversion will be adopted by organisations wishing to describe their sector’s failings, and the adoption will do no work (§3.6).

The trust argument does not transfer at all, because it does not obtain in the first place (§5.14). An organisation adopting epistocracy to demonstrate its accountability has adopted the ritual and not the mechanism, and Power’s analysis predicts exactly this migration⁹⁰: the audit form travels, and the capacity to withhold assurance does not travel with it.

What does transfer is a single structural property and its conditions: **the capacity to reopen a conclusion cannot be controlled by the authority whose conclusion would be reopened**, secured by independent appointment, standing for the sealed party, non-discretionary thresholds, recorded refusals, retaliation protection, and periodic review against later evidence. That property is indifferent to sector. It asks nothing about the goodness of the people involved, which is why it survives contact with them.

9. Conclusion

Universities have long positioned themselves as custodians of knowledge. This paper has argued that in respect of their evaluative powers they should be constituted as fiduciaries, owing duties of loyalty, care and answerability to those governed by the standards they apply — students, applicants, junior colleagues, rejected researchers — who bear the costs of those standards and cannot monitor or replicate the deliberation that produced them. That is a proposal for how the university ought to be constituted, not a description of a duty English law presently imposes; the distinction is drawn at §1.1 and is load-bearing. These obligations cannot be discharged under prevailing regimes of managerial optics, academic capitalism and superficial compliance.

Nor can they be discharged by a more candid version of the same regimes. The diagnosis of Chapters 2 and 3 identifies pathologies that do not require anyone to intend them, and a remedy addressed to intention therefore misses. The remedy the sector reaches for is of exactly that kind: audit boards that will look carefully, report honestly, and thereby restore trust. Under the seal, however careful, the looking cannot reach the residue of the standard under which the looking is done; under Power’s

⁹⁰ Power, *The Audit Society*.

analysis, the trust-restoring function destroys the withholding capacity that would make the looking worth anything. The prescription is the disease.

What replaces it is smaller, harder, and does not flatter anyone. Standards must be capable of losing. The sealed party must hold a key that the institution cannot take back. The body that reviews must not be the body that appoints, resources or dissolves it. Reasons must be recorded, not to explain the standard in full — no account reaches the whole of it — but to fix the claim before the outcome arrives, so that the outcome can contradict it. Review dates must be set in advance and must examine the excluded, because the excluded are where a standard's failure is stored. And no part of this may be justified by the trust it produces, because a mechanism that must produce trust has lost the capacity to withhold it, and a mechanism that cannot withhold is decoration with a budget.

Epistemocracy so understood reclaims the university's democratic and epistemic missions on a narrow footing, and a firm one. It aligns institutional governance with democratic epistemology⁹¹, fiduciary ethics⁹², and the tradition of intellectual virtue⁹³ — while insisting that virtue is what the mechanism protects, not what protects the mechanism. It extends critical pedagogy⁹⁴ into governance by taking seriously the one question those traditions pose most sharply: to whom is this institution answerable, and what can they *do*?

Nothing here promises that universities adopting these mechanisms will be better regarded. It promises that their standards will be capable of losing. Those are different goods, and an institution that finds itself preferring the first has located its own hollowing more precisely than any audit board could — which is, at least, a finding.

Universities must be more than sites of credential production or brand management. But the demand cannot be met by exhortation, because the institution exhorted cannot see the residue it would need to correct. It can be met by architecture: by handing the key to the person the standard excluded, and by fixing, in advance, the date on which the institution does not get to decide what the evidence says.

⁹¹ Anderson, 'The Epistemology of Democracy'.

⁹² Frankel, *Fiduciary Law*; Smith, *The Law of Loyalty*; Miller, 'The Identification of Fiduciary Relationships'.

⁹³ Zagzebski, *Virtues of the Mind*.

⁹⁴ Freire, *Pedagogy of the Oppressed*; Darder, *Reinventing Paulo Freire*.

Version History

Ver	Description	Epistemic Impact	DOI	Date
1.0	Initial release, under the title 'Epistemocracy in Higher Education: A Proposal for Fiduciary and Epistemic Accountability in the University' (retained at v2.0; the paper was retitled 'Standards That Can Lose: An Epistemic-Fiduciary Constitution for the University' at v3.0).	None (baseline release).	None issued.	2025-06-19
2.0	Substantial structural expansion; new chapters on Naming as Epistemic Tool and Broader Intellectual and Societal Implications; integration of neologisms (optocratic drift, fiducial hollowing, epistemic inversion); enhanced comparative analysis across universities, corporations, NGOs, and democratic institutions; updated citation style to APA; streamlined earlier chapters for concision and clarity.	These revisions sharpen the conceptual vocabulary for diagnosing epistemic distortions, embed fiduciary-epistemic accountability more firmly in institutional governance, and extend the reach of epistemocracy as a transferable societal model. The 2nd edition transforms the work from a critical-analytical essay into a comprehensive institutional blueprint, significantly increasing its theoretical rigour and practical applicability.	10.5281/zenodo.17293583	2025-09-15
3.0	Substantive corrigendum to the remedy; the diagnosis (Chs 2-3) is retained. Fiduciary claim recast as a proposal – universities <i>should be constituted as</i> epistemic fiduciaries in respect of specified evaluative powers, not that they are so in positive law (status follows an undertaking of loyalty per <i>Mothew</i> ; the standing gap evidenced by the academic-judgment exclusion in HEA 2004 s 12, which is treated as distinct from but overlapping the evaluative standard). Draws on the companion diagnostic study, 'The Fiduciary Breach of the Modern University', adopting its vocabulary of <i>fiduciary breach</i> and <i>epistemic capture</i> under an expressly narrowed construction: breach as the failure of answerable defeasibility rather than of candour to the public, and capture named at the level of a single evaluative power (§§1.1, 3.1, 5.12). Seal narrowed from the absolute claim to the defensible one (no authority can <i>exhaustively and self-validatingly</i> represent its own standard); transparency recast as necessary-but-insufficient, and periodic review's primacy re-derived from the unreach residue together with clientelism and evaluative-layer capping. Remedy rebuilt : 'fiduciary transparency' → fiduciary answerability; periodic review against later evidence made primary and given a three-tier evidential rule (trigger / burden-shift / defeat above a pre-specified false-negative rate, base rates restored, with the comparator-contamination limit stated and matched by a defeat condition at §6.5); standing to <i>compel</i> review vested in affected parties on a bounded, reviewable threshold with retaliation protection; appointing authority specified and non-self-control required, with independence disaggregated into five axes; plurality demoted; the legitimacy framing struck and inverted. Added : a law of consequences for a defeated standard and a standard-attribution provision routing findings to the constitutor (§§5.12-5.13); a data-protection analysis under the UK GDPR and DPA 2018 as amended by the Data (Use and Access) Act 2025, corrected on four points (§5.2); a terminological note distinguishing the proposal from Estlund's epistocracy and Taleb's epistemocracy (§1.1); and the empirical warrant for the diagnosis (Espeland & Sauder, Wilsdon, Strathern). Retitled from 'Epistemocracy in Higher Education: A Proposal for Fiduciary and Epistemic Accountability in the University' to 'Standards That Can Lose: An Epistemic-Fiduciary Constitution for the University', leading with the doctrinal contribution and demoting 'epistemocracy' from title to the name of the model developed within (abstract, §§1.1, 4). Series strand reassigned from higher-education governance to fiduciary and public-law governance; keywords reordered to lead with the fiduciary and public-law contribution. Restructured : Ch 1 recast as a standard introduction; principles and mechanisms separated into Chs 4 and 5; duplicative chapters removed; sectioning limited to two levels; citations converted to OSCOLA. The text no longer argues against its own earlier versions – this row is where the change of position is recorded.	Reverses the direction of the remedy. Versions 1.0-2.0 diagnosed fiducial hollowing and optocratic drift and then proposed an oversight architecture that reproduced both: bodies no beneficiary could invoke, appointed on the face of the text by the audited party, defended by plurality rather than independence, and justified by their capacity to produce legitimacy – which is Power's <i>Audit Society</i> thesis, deployed in Chapter 3 against everyone else and walked into in Chapter 4. The v3 remedy rests on exposure rather than representation: a standard must remain capable of losing by a route its holder does not control. Three things are consequently withdrawn – transparency as a remedy, virtue as a check, and trust as an objective – and the model claims materially less than v2 while becoming, for the first time, testable. The paper now states the conditions under which it is itself defeated, and predicts its own most likely mode of misadoption. Users of v2's §4.4-4.5 as an implementation template should treat that text as withdrawn.	10.13140/RG.2.2.14076.09603 10.5281/zenodo.21468283	2026-07-18

Table of Cases

Bristol and West Building Society v Mothew [1998] Ch 1

Keech v Sandford (1726) Sel Cas Ch 61

Rukhadze v Recovery Partners GP Ltd [2025] UKSC 10

Table of Legislation

Companies Act 2006

Data (Use and Access) Act 2025

Data Protection Act 2018

Higher Education Act 2004

Retained Regulation (EU) 2016/679 (UK General Data Protection Regulation)

Bibliography

Anderson E, 'The Epistemology of Democracy' (2006) 3 *Episteme* 8

Barnett R, *Realizing the University in an Age of Supercomplexity* (SRHE/Open University Press 2000)

—— *The Philosophy of Higher Education: A Critical Introduction* (Routledge 2022)

—— *Realizing the Ecological University: Eight Ecosystems, Their Antagonisms and a Manifesto* (Bloomsbury 2025)

Biesta G, *The Beautiful Risk of Education* (Routledge 2013)

Darder A, *Reinventing Paulo Freire: A Pedagogy of Love* (2nd edn, Routledge 2017)

Dotson K, 'A Cautionary Tale: On Limiting Epistemic Oppression' (2012) 33 *Frontiers: A Journal of Women Studies* 24

—— 'Conceptualizing Epistemic Oppression' (2014) 28 *Social Epistemology* 115

DeMott DA, 'The Fiduciary Character of Agency and the Interpretation of Instructions' in Andrew S Gold and Paul B Miller (eds), *Philosophical Foundations of Fiduciary Law* (OUP 2014) 321

Espeland WN and Sauder M, 'Rankings and Reactivity: How Public Measures Recreate Social Worlds' (2007) 113 *American Journal of Sociology* 1

Estlund DM, *Democratic Authority: A Philosophical Framework* (Princeton University Press 2008)

Foucault M, *Discipline and Punish: The Birth of the Prison* (Alan Sheridan tr, Pantheon 1977)

Frankel T, *Fiduciary Law* (OUP 2011)

Freire P, *Pedagogy of the Oppressed* (Myra Bergman Ramos tr, Continuum 1970)

Fricke M, *Epistemic Injustice: Power and the Ethics of Knowing* (OUP 2007)

Kahl P, 'Conformance Is Not Constitution: Provenance, Answerability, and Standing' (Lex et Ratio Ltd 2026)
<<https://doi.org/10.13140/RG.2.2.28710.28486>>

- ‘Delegated Discretion: When Authority Lacks Evaluative Control in Fiduciary and Administrative Law’ (version 2.0, Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20920750>>
- ‘Directors’ Epistemic Duties and Fiduciary Openness: A Cross-Cultural and Interdisciplinary Framework for Corporate Governance’ (2nd edn, Lex et Ratio Ltd 2025) <<https://doi.org/10.13140/RG.2.2.28155.17449>>
- ‘Epistemic Justice and Institutional Responsibility in Academia: Toward a Comprehensive Framework for Epistemic Justice in Higher Education’ (2nd edn, Lex et Ratio Ltd 2025) <<https://doi.org/10.13140/RG.2.2.11063.38563>>
- ‘Answerable Knowledge: Procedural and Epistemic Authority in the Research Organisation’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20439758>>
- ‘Epistemic Humility as Fiduciary Obligation: Entrusted Discretion and Responsibility for Belief’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.18440605>>
- ‘The Fiduciary Breach of the Modern University: Fiduciary Law and the Epistemic Constitution of the Fourth Estate’ (Lex et Ratio Ltd 2025) <<https://doi.org/10.5281/zenodo.17618878>>
- ‘The Inside of Generatedness: Representational Sealing and the Materials of Temporal Order’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20353457>>
- ‘Out-of-Distribution Failure as a Structural Signal of Epistemic Enclosure’ (Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.19353400>>
- ‘The Recognition Game: Metric Governance and the Stabilisation of Epistemic Clientelism’ (version 2.0, Lex et Ratio Ltd 2026) <<https://doi.org/10.5281/zenodo.20774711>>
- Medina J, *The Epistemology of Resistance: Gender and Racial Oppression, Epistemic Injustice, and Resistant Imaginations* (OUP 2013)
- Miller PB, ‘The Fiduciary Relationship’ in Andrew S Gold and Paul B Miller (eds), *Philosophical Foundations of Fiduciary Law* (OUP 2014) 63
- ‘The Identification of Fiduciary Relationships’ in Evan J Criddle, Paul B Miller and Robert H Sitkoff (eds), *The Oxford Handbook of Fiduciary Law* (OUP 2019) 366
- Mouffe C, *Agonistics: Thinking the World Politically* (Verso 2013)
- Office of the Independent Adjudicator for Higher Education, *Guidance on the Rules* <<https://www.oiahe.org.uk/about-us/our-scheme/our-rules/guidance-on-the-rules/rule-4-and-rule-5/>> accessed 17 July 2026
- Power M, *The Audit Society: Rituals of Verification* (OUP 1997)
- Slaughter S and Rhoades G, *Academic Capitalism and the New Economy: Markets, State, and Higher Education* (Johns Hopkins University Press 2004)
- Smith L, *The Law of Loyalty* (OUP 2023)
- Strathern M (ed), *Audit Cultures: Anthropological Studies in Accountability, Ethics and the Academy* (Routledge 2000)
- Taleb NN, *The Black Swan: The Impact of the Highly Improbable* (Random House 2007)
- Wilsdon J and others, *The Metric Tide: Report of the Independent Review of the Role of Metrics in Research Assessment and Management* (HEFCE 2015)
- Zagzebski LT, *Virtues of the Mind: An Inquiry into the Nature of Virtue and the Ethical Foundations of Knowledge* (CUP 1996)